NATURAL LANGUAGE EVALUATION WITH HUMANS IN THE
LOOP AND STATISTICAL ESTIMATORS

A DISSERTATION
SUBMITTED TO THE DEPARTMENT OF COMPUTER SCIENCE
AND THE COMMITTEE ON GRADUATE STUDIES
OF STANFORD UNIVERSITY
IN PARTIAL FULFILLMENT OF THE REQUIREMENTS
FOR THE DEGREE OF
DOCTOR OF PHILOSOPHY

Arun Tejasvi Chaganty
June 2018

I certify that I have read this dissertation and that, in my opinion, it is fully adequate in scope and quality as a dissertation for the degree of Doctor of Philosophy.

_____

(Percy S. Liang)    Principal Adviser

I certify that I have read this dissertation and that, in my opinion, it is fully adequate in scope and quality as a dissertation for the degree of Doctor of Philosophy.

_____

(Christopher D. Manning)

I certify that I have read this dissertation and that, in my opinion, it is fully adequate in scope and quality as a dissertation for the degree of Doctor of Philosophy.

_____

(Michael S. Bernstein)

Approved for the Stanford University Committee on Graduate Studies

_____

# Preface

In natural language tasks such as knowledge base population, text summarization or open-response question answering, a significant challenge is simply evaluating the performance of automated systems because of the large diversity of possible outputs. Existing fully-automatic methods for evaluating these systems rely on an incomplete set of annotated references which lead to systematic biases against certain system improvements: in other words, genuinely good ideas are systematically discarded simply because of limitations in our evaluation methodology. As a result, human evaluation, which can be prohibitively expensive, has remained the de-facto mode of evaluation for these tasks. In this work, we show how one can decrease the costs of incorporating human feedback through the design of appropriate statistical estimators.

First, we consider the setting where the output produced by systems is "dense", in that we expect significant overlap between the output of different systems. Naively combining annotations from different systems leads to a representation bias. Here, we show that cost of obtaining human feedback can be significantly amortized by using a novel importance-reweighted estimator. We apply this estimator to design a new evaluation methodology for knowledge base population and show that cost of evaluating precision and recall within this framework can be reduced by a factor of 4.

Next, we consider the "sparse" setting wherein few, if any, systems ever produce identical output. Traditionally, the community has relied on similarity-based automatic metrics such as BLEU or ROUGE to compare the outputs produced by different systems. Unfortunately, these metrics have been shown to poorly correlate with human judgment and thus introduce bias in evaluation. We derive an unbiased estimator that optimally combines these automatic metrics with human feedback. Our theoretical results allow us to

characterize potential cost reductions only in terms of the tasks' subjectivity, measured by inter-annotator variance, and the automatic metrics' quality, measured by correlation with human judgments. On two popular natural language generation tasks, question answering and summarization, we empirically show that currently we can achieve at most a 7–13% reduction in cost on two tasks, exposing fundamental limitations in debiasing current automatic metrics.

Finally, we show how human feedback can be incorporated into systems in real-time to deploy high-accuracy systems starting with zero training examples. We build systems that learn "on-the-job" by using human feedback to resolve uncertainty until it is confident in its predictions. Our key statistical idea here is to cast the problem as a stochastic game based on Bayesian decision theory, which allows us to balance latency, cost, and accuracy objectives in a principled way. When tested on three classification tasks—named-entity recognition, sentiment classification, and image classification— we obtained an order of magnitude reduction in cost compared to full human annotation, while also boosting performance relative to the expert provided labels.

# Acknowledgments

I would like to thank. . .

# Contents

# List of Tables

# List of Figures

# Chapter 1

# Introduction

One of the most exciting applications of natural language processing ahead of us is the development of systems that will be able to *summarize information* for us. For example, we will one day be able to rely on text summarization (TS) systems to condense the most salient information from one or more related articles into as few words as necessary (Figure 1.1a). Knowledge base population (KBP) systems will be able to read large document troves (or the Internet) and present succinct, structured, summaries about the people and organizations mentioned within (Figure 1.1b). Finally, we will be be able to get targeted summaries through open-ended question answering systems (OQA) (Figure 1.1c).

Indeed, this broad vision of summarization has been the goal of many natural language processing systems dating as far back as the 1950s (Luhn, 1958). Research has shown that having access to document summaries significantly improves user satisfaction and their ability to complete fact-gathering tasks (Mani et al., 1999; Mckeown et al., 2005). Building these systems, however, remains a challenge. Today, we are seeing a resurgence in interest in automatic summarization systems, with over 150 papers published at top NLP conferences in just the last year (2017–18).

Still, despite years of effort, it is unclear whether we are actually making forward progress on these tasks. Brandow et al. (1995) conducted a large scale human evaluation of different summarization systems and found that the baseline of simply taking the lead-paragraph significantly outperformed automatic systems. Ten years later, Passonneau et al. (2005) report that half of the systems participating in the DUC-2005 summarization
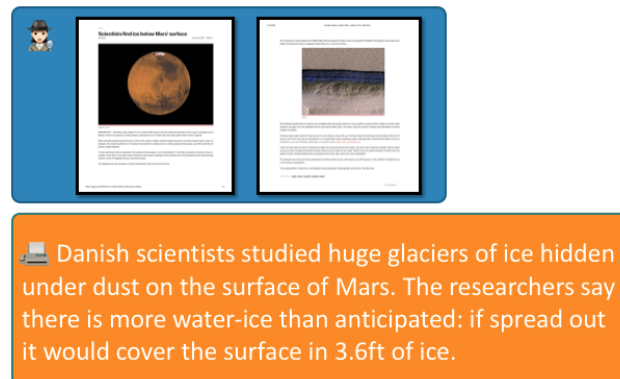
(a)  Text summarization.



(b)  Knowledge base population.
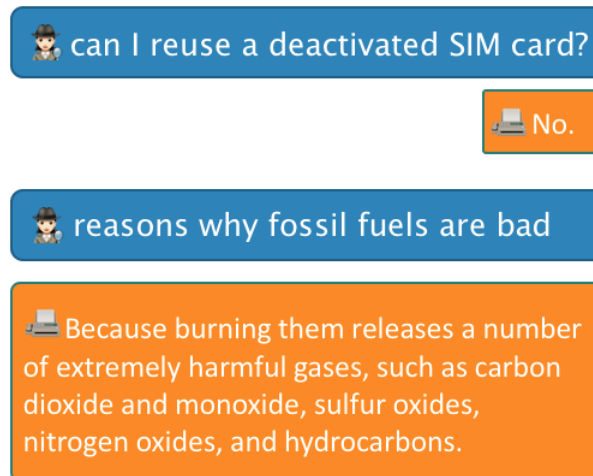


(c)  Open-response question answering.

Figure 1.1:  An overview of some information summarization tasks. (a) Text summariza-
tion: seeks to identify the most salient information contained within one (or more) article(s)
and describe this information in as few words as necessary. (b) Knowledge base popula-
tion: reads documents from a large corpus and generates linked entity-centric summaries
for every person and organization mentioned within. (c) Open-response question answer-
ing: provides very-targeted summaries that address a specific question.

Figure 1.2: Much of our evaluation methodology relies on a representative evaluation dataset that contains paired inputs and outputs. In this work, we argue that this evaluation methodology relies on the outputs belonging to a small closed class of labels (e.g. positive or negative sentiment) to be able to easily judge output of system. In the summarization tasks we are interested in, this assumption does not hold. We consider the data to be *incomplete*: .

challenge did worse than the baseline. Even today, we find that recent "state-of-the-art" neural network based summarization systems fare poorly on human evaluations of language quality relative to simpler extractive systems. While there are many reasons for thsi slow progress, we believe that the lack of an effective evaluation methodology remains one of the most important ones.

The ultimate goal of this thesis is to enable the development of better information summarization systems by addressing systemic problems in how we evaluate them. In this work, we attribute the inherent *incompleteness* of existing evaluation datasets as the key bottleneck to reliably measuring the performance of summarization systems. We posit that collecting on-demand human feedback is fundamental to obtaining meaningful measures of system quality by resolving this incompleteness. Our key technical contribution is applying techniques from statistical estimation to reliably integrate human feedback in a cost-effective manner.

## 1.1 The incompleteness of static evaluation sets

At its core, the prevalent evaluation methodology today relies on a evaluation dataset that contains pairs of query inputs and expected outputs (Figure **??**). On classification problems, e.g. sentiment classification or topic identification, the output typically belongs to a small closed class (e.g. positive, neutral or negative sentiment). When a system also predicts an output belonging to this closed class, it is very easy to say whether or not the system made a mistake. Thus, in these settings, we can measure the quality of different systems on the classification task by simply collecting a sufficiently large evaluation dataset.[1]

---

[1]Of course, care must be taken to ensure the evaluation dataset collected is reflective of the end-goal for the task.

Unfortunately, in the summarization tasks we've seen above, the desired output is not simply a class label but rather an arbitrary piece of text (e.g. for TS and OQA) or an arbitrarily large collection of facts (e.g. for KBP). The assumption that there exists a unique knowable correct output simply does not hold. In this sense, the dataset is *incomplete*: it either does not contain every possible correct answer or element. Let's look at some examples of how incompleteness manifests in practice and what its consequences on evaluation could be.

Figure 1.4a shows the candidate output of two KBP systems for Carrie Fisher, the late actress who played Princess Leia in the Star Wars movie franchise. The incomplete KBP reference data only specifies that Ms. Fisher is an actress and thus the evaluation is unable to judge whether the system that identifies her to also be an author (which is correct) is better or worse than one that identifies her to also be a princess (which is incorrect): by default, the evaluation penalizes both decisions equally. As a consequence, using this evaluation as an empirical guide leads researchers to avoid improvements that would identify Ms. Fisher as an actress.

Next, consider an example in question answering: here, the fundamental problem is that there are many ways to express the same answer (Figure 1.4b). It is not yet easy to automatically identify whether two phrases mean the same thing and hence fairly judge the output systems produce. More subtly, we show that this evaluation is biased towards easier questions with short answers that are more likely to be reproduced by a system.

Finally, consider the evaluation of text summarization: a system generated summary will never exactly match the one in the evaluation dataset. A common practice in the community is to use a word-overlap based similarity score such as BLEU (Papineni et al., 2002) or ROUGE (Lin and Rey, 2004). Unfortunately, these automatic metrics have been shown to correlate extremely poorly with human judgment (Novikova et al., 2017). Figure 1.4c shows an example of a published system that learns to game this metric by appending seemingly random words at the end of a summary.
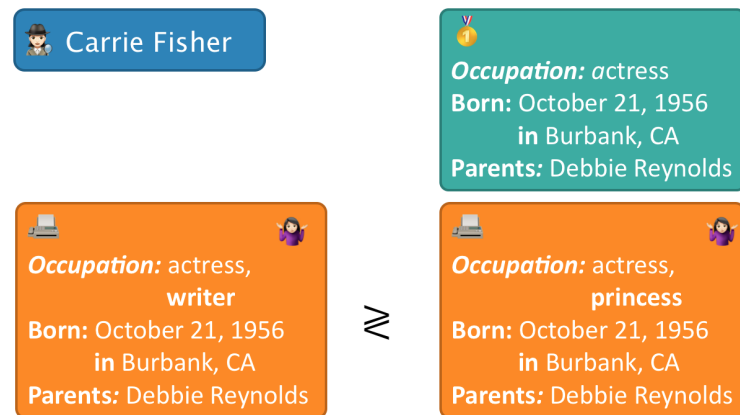
(a) Knowledge base population.



(b) Open-response question answering.
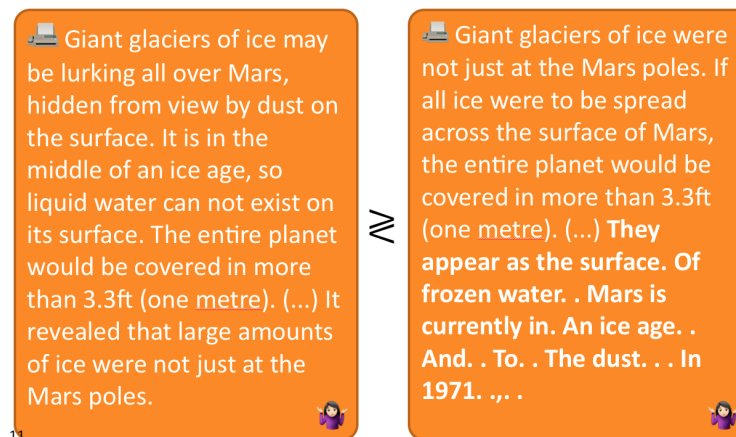


(c) Text summarization.

Figure 1.3: Examples where the reference data is unable to evaluate a system's response: (a) Knowledge base population: here, the incomplete reference data is unable to identify that it is a true that Carrie Fisher was also an author, but not that she is a 'princess'. (b) Open-response question answering: it is common for systems to output a paraphrase of the reference answer and thus to be judged incorrect despite reporting the right answer. (c) Text summarization: systems rarely produce output that is identical to the reference answer and word-overlap based similarity measures often prefer bad responses with high lexical similarity over good responses that more more lexically distinct.

## 1.2 Addressing incompleteness with human feedback

We've just seen how the incompleteness of our evaluation sets can introduce bias. How fundamental is incompleteness? In the setting of knowledge base population, it is possible to minimize the impact of incompleteness by exhaustively annotating a sufficiently large document collection. We estimate that this would cost at least $1 million; of course, any adjustment in the task definition (e.g. the relation schema) would require a fresh batch of annotations. In this sense, handling incompleteness in KBP evaluation can be viewed as a prudent cost-saving measure. On the other hand, the number of correct wordings for an answer in text summarization or open-ended question answering are nearly endless; we doubt that any amount of money would be sufficient to ameliorate the problem here.

The crux of the problem is estimating the impact of instances that are unknown to the automatic evaluation that relies on a static dataset. We exploit the fact that the answers to these instances are obvious to humans and propose asking people for feedback *on-demand* using advances in crowdsourcing. In this work, we advocate moving from annotating data in batches to annotating data *on-demand*.

Apart from fixing the problem of incompleteness, incorporating human feedback has several ancillary benefits. First, it ensures that our evaluation does not diverge from the end goal. Second, it provides *qualitative* guidance to help us identify opportunities for improvements. As an example of this latter point, in Chapter 3 we show how combining human labels with statistical estimators can be used to conduct fine-grained error analysis for KBP, and in Chapter 4 we show how collecting human edits as feedback can be used for error analysis in text generation.

## 1.3 Integrating human feedback with statistical estimators

While human evaluation is often regarded as a gold standard for evaluation, it is also considered to be too expensive to be used as when iterating on models. The key question this thesis tries to answer then is: can we reduce the cost of human annotation in evaluation

while maintaining its fidelity? Our core contribution is in addressing the problem of incompleteness for two extremes: problems where incompleteness is large but finite (KBP), and problems where incompleteness is truly infinite (TS and OQA). In both settings, we hold *unbiased* estimation as our bar: the methods we describe are guaranteed to produce the same results in expectation as would human evaluation.

**Amortizing costs when incompleteness is finite.**    In settings where incompleteness is finite, it is likely that two different systems produce an overlapping set of outputs, suggesting that we may can leverage human annotations obtained for one system to evaluate another. The key challenge is guarding against representation bias: we don't want to skew the evaluation towards a particular class of systems simply because our evaluation data came from these systems. We tackle this problem in Chapter 3 using a novel importance-reweighted estimator. We apply the estimator to evaluate KBP systems and show that we are able to reduce the cost of obtaining human annotations by a factor of 4.

**Finding limitations when incompleteness is infinite.**    On the other end of the spectrum, in tasks like text summarization or open-ended question answering, the output produced consists of free form text: it is extremely unlikely that two systems will ever agree on the output they produce. Here, it seems natural to rely on some "similarity" measure that may allow us to match two similar, but non-identical responses. In Chapter 4, we derive an optimal estimator to combine such a similarity metric with human feedback based on control variates (Owen, 2013). Our theoretical analysis allows us to characterize when it is possible to reduce human annotation costs while guaranteeing unbiasedness. We show that for both text summarization and open-ended question answering current cost savings are modest, about 10–15%, owing to both the poor quality of existing automatic metrics and the inherent annotator variance for these subjective tasks.

**Beyond evaluation.**    Thus far, we have focused on the use of human feedback during evaluation. In the Chapter 5, we show how human feedback can be effectively integrated *at test-time*. We consider an "on-the-job" setting, where our model learns as inputs arrive: we use real-time crowdsourcing to resolve uncertainty where needed and output our prediction

once the model is confident.  As the model improves over time, the reliance on crowd-sourcing queries decreases.  We cast our setting as a stochastic game based on Bayesian decision theory, which allows us to balance latency, cost, and accuracy objectives in a principled way.  Computing the optimal policy is intractable, so we develop an approximation based on Monte Carlo Tree Search.  We tested our approach on three datasets—named-entity recognition, sentiment classification, and image classification.  On the NER task we obtained more than an order of magnitude reduction in cost compared to full human annotation, while boosting performance relative to the expert provided labels.

# Chapter 2

# Setup

placeholder.

Newsblaster Even with these simple systems, Mckeown et al. (2005) showed that providing users such multi-document summaries help.

Only recently have abstractive systems based on. Multi-document summarization SUMMONS, NewsBlaster

Topic driven summarization

Information extraction based summarization.

Macro-rule systems.

# Chapter 3

# Importance-reweighted estimators for unbiased on-demand evaluation of knowledge base population

Knowledge base population (KBP) systems take in a large document corpus and extract entities and their relations. Thus far, KBP evaluation has relied on judgements on the pooled predictions of existing systems. We show that this evaluation is problematic: when a new system predicts a previously unseen relation, it is penalized even if it is correct. This leads to significant bias against new systems, which counterproductively discourages innovation in the field. Our first contribution is a new importance-sampling based evaluation which corrects for this bias by annotating a new system's predictions on-demand via crowdsourcing. We show this eliminates bias and reduces variance using data from the 2015 TAC KBP task. Our second contribution is an implementation of our method made publicly available as an online KBP evaluation service. We pilot the service by testing diverse state-of-the-art systems on the TAC KBP 2016 corpus and obtain accurate scores in a cost effective manner.

## 3.1 Introduction

Harnessing the wealth of information present in unstructured text online has been a long standing goal for the natural language processing community. In particular, knowledge
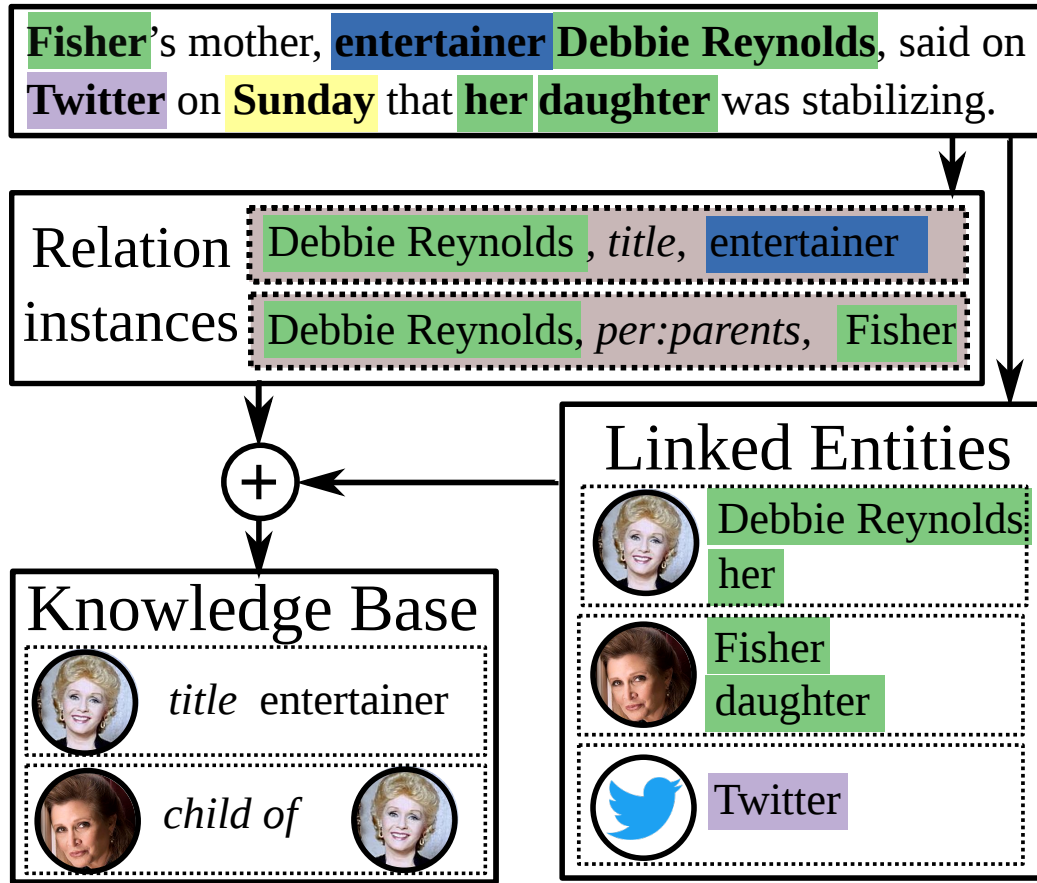
Figure 3.1: An example describing entities and relations in knowledge base population.

base population seeks to automatically construct a knowledge base consisting of relations between entities from a document corpus. Knowledge bases have found many applications including question answering (Berant et al., 2013; Fader et al., 2014; Reddy et al., 2014), automated reasoning (Kalyanpur et al., 2012) and dialogue (Han et al., 2015).

Evaluating these systems remains a challenge as it is not economically feasible to exhaustively annotate every possible candidate relation from a sufficiently large corpus. As a result, a pooling-based methodology is used in practice to construct datasets, similar to them methodology used in information retrieval (Jones and Rijsbergen, 1975; Harman, 1993). For instance, at the annual NIST TAC KBP evaluation, all relations predicted by participating systems are pooled together, annotated and released as a dataset for researchers to develop and evaluate their systems on. However, during development, if a new system

predicts a previously unseen relation it is considered to be wrong even if it is correct. The discrepancy between a system's true score and the score on the pooled dataset is called *pooling bias* and is typically assumed to be insignificant in practice (Zobel, 1998).

The key finding of this paper contradicts this assumption and shows that the pooling bias is actually significant, and it penalizes newly developed systems by 2% $F_1$ on average (Section 3.3). Novel improvements, which typically increase scores by less than 1% $F_1$ on existing datasets, are therefore likely to be clouded by pooling bias during development. Worse, the bias is larger for a system which predicts qualitatively different relations systematically missing from the pool. Of course, systems participating in the TAC KBP evaluation do not suffer from pooling bias, but this requires researchers to wait a year to get credible feedback on new ideas.

This bias is particularly counterproductive for machine learning methods as they are trained assuming the pool is the complete set of positives. Predicting unseen relations and learning novel patterns is penalized. The net effect is that researchers are discouraged from developing innovative approaches, in particular from applying machine learning, thereby slowing progress on the task.

Our second contribution, described in Section 4.3, addresses this bias through a new evaluation methodology, *on-demand evaluation*, which avoids pooling bias by querying crowdworkers, while minimizing cost by leveraging previous systems' predictions when possible. We then compute the new system's score based on the predictions of past systems using importance weighting. As more systems are evaluated, the marginal cost of evaluating a new system decreases. We show how the on-demand evaluation methodology can be applied to knowledge base population in Section 3.5. Through a simulated experiment on evaluation data released through the TAC KBP 2015 Slot Validation track, we show that we are able to obtain unbiased estimates of a new systems score's while significantly reducing variance.

Finally, our third contribution is an implementation of our framework as a publicly available evaluation service at https://kbpo.stanford.edu, where researchers can have their own KBP systems evaluated. The data collected through the evaluation process could even be valuable for relation extraction, entity linking and coreference, and will also be made publicly available through the website. We evaluate three systems on the 2016

$$\text{Humans} \dashleftarrow \dashrightarrow i_1 : (s_1, r, o_1, p_1) \checkmark$$

$$i_2 : (s_1, r, o_2, p_2) \checkmark$$

$$\text{System A} \qquad i_3 : (s_1, r, o_3, p_3) \checkmark$$

$$\text{System B} \qquad i_4 : (s_1, r, o_2, p_4) \checkmark$$

$$i_5 : (s_1, r, o_3, p_5) \times$$

$$\text{System C} \qquad i_6 : (s_1, r, o_4, p_6) \checkmark$$

Figure 3.2: In pooled evaluation, an evaluation dataset is constructed by labeling relation instances collected from the pooled systems (A and B) and from a team of human annotators (Humans). However, when a new system (C) is evaluated on this dataset, some of its predictions ($i_6$) are missing and can not be fairly evaluated. Here, the precision and recall for C should be $\frac{3}{3}$ and $\frac{3}{4}$ respectively, but its evaluation scores are estimated to be $\frac{2}{3}$ and $\frac{2}{3}$. The discrepancy between these two scores is called *pooling bias*.

TAC KBP corpus for about \$150 each (a fraction of the cost of official evaluation). We believe the public availability of this service will speed the pace of progress in developing KBP systems.

## 3.2 Background

In knowledge base population, each relation is a triple (SUBJECT, PREDICATE, OBJECT) where SUBJECT and OBJECT are some globally unique entity identifiers (e.g. Wikipedia page titles) and PREDICATE belongm to a specified schema.[1] A KBP system returns an output in the form of *relation instances* (SUBJECT, PREDICATE, OBJECT, PROVENANCE),

---

[1] The TAC KBP guidelines specify a total of 65 predicates (including inverses) such as `per:title` or `org:founded_on`, etc. Subject entities can be people, organizations, geopolitical entities, while object entities also include dates, numbers and arbitrary string-values like job titles.

where PROVENANCE is a description of where exactly in the document corpus the relation was found. In the example shown in Figure 3.1, CARRIE FISHER and DEBBIE REYNOLDS are identified as the subject and object, respectively, of the predicate CHILD OF, and the whole sentence is provided as provenance. The provenance also identifies that CARRIE FISHER is referenced by **Fisher** within the sentence. Note that the same relation can be expressed in multiple sentences across the document corpus; each of these is a different relation instance.

**Pooled evaluation.** The primary source of evaluation data for KBP comes from the annual TAC KBP competition organized by NIST (Ji et al., 2011). Let $E$ be a held-out set of *evaluation entities*. There are two steps performed in parallel: First, each participating system is run on the document corpus to produce a set of relation instances; those whose subjects are in $E$ are labeled as either positive or negative by annotators. Second, a team of annotators identify and label correct relation instances for the evaluation entities $E$ by manually searching the document corpus within a time budget (Ellis et al., 2012). These labeled relation instances from the two steps are combined and released as the evaluation dataset. In the example in Figure 3.2, systems A and B were used in constructing the pooling dataset, and there are 3 distinct relations in the dataset, between $s_1$ and $o_1$, $o_2$, $o_3$.

A system is evaluated on the precision of its predicted relation instances for the evaluation entities $E$ and on the recall of the corresponding predicted *relations* (not instances) for the same entities (see Figure 3.2 for a worked example). When using the evaluation data during system development, it is common practice to use the more lenient `anydoc` score that ignores the provenance when checking if a relation instance is true. Under this metric, predicting the relation (CARRIE FISHER, CHILD OF, DEBBIE REYNOLDS) from an ambiguous provenance like "**Carrie Fisher** and **Debbie Reynolds** arrived together at the awards show" would be considered correct even though it would be marked wrong under the official metric.

## 3.3 Measuring pooling bias

The example in Figure 3.2 makes it apparent that pooling-based evaluation can introduce a systematic bias against unpooled systems. However, it has been assumed that the bias is insignificant in practice given the large number of systems pooled in the TAC KBP evaluation. We will now show that the assumption is not valid using data from the TAC KBP 2015 evaluation.[2]

**Measuring bias.**    In total, there are 70 system submissions from 18 teams for 317 evaluation entities ($E$) and the evaluation set consists of 11,008 labeled relation instances.[3] The original evaluation dataset gives us a good measure of the true scores for the participating systems. Similar to Zobel (1998), which studied pooling bias in information retrieval, we simulate the condition of a team not being part of the pooling process by removing any predictions that are unique to its systems from the evaluation dataset. The pooling bias is then the difference between the true and unpooled scores.

**Results.**    Figure 3.3 shows the results of measuring pooling bias on the TAC KBP 2015 evaluation on the $F_1$ metric using the official and `anydoc` scores.[45] We observe that even with lenient `anydoc` heuristic, the median bias (2.05% $F_1$) is much larger than largest difference between adjacently ranked systems (1.5% $F_1$). This experiment shows that pooling evaluation is significantly and systematically biased against systems that make novel predictions!

---

[2]Our results are not qualitatively different on data from previous years of the shared task.

[3]The evaluation set is actually constructed from compositional queries like, "what does Carrie Fisher's parents do?": these queries select relation instances that answer the question "who are Carrie Fisher's parents?", and then use those answers (e.g. "Debbie Reynolds") to select relation instances that answer "what does Debbie Reynolds do?". We only consider instances selected in the first part of this process.

[4]We note that `anydoc` scores are on average $0.88\% F_1$ larger than the official scores.

[5] The outlier at rank 36 corresponds to a University of Texas, Austin system that only filtered predictions from other systems and hence has no unique predictions itself.

## 3.4 On-demand evaluation with importance sampling

Pooling bias is fundamentally a sampling bias problem where relation instances from new systems are underrepresented in the evaluation dataset. We could of course sidestep the problem by exhaustively annotating the entire document corpus, by annotating all mentions of entities and checking relations between all pairs of mentions. However, that would be a laborious and prohibitively expensive task: using the interfaces we've developed (Section 4.5), it costs about $15 to annotate a single document by non-expert crowdworkers, resulting in an estimated cost of at least $1,350,000 for a reasonably large corpus of 90,000 documents (Dang, 2016). The annotation effort would cost significantly more with expert annotators. In contrast, *labeling* relation instances from system predictions can be an order of magnitude cheaper than finding them in documents: using our interfaces, it costs only about $0.18 to verify each relation instance compared to $1.60 per instance extracted through exhaustive annotations.

We propose a new paradigm called on-demand evaluation which takes a lazy approach to dataset construction by annotating predictions from systems *only when they are underrepresented*, thus correcting for pooling bias as it arises. In this section, we'll formalize the problem solved by on-demand evaluation independent of KBP and describe a cost-effective solution that allows us to accurately estimate evaluation scores without bias using importance sampling. We'll then instantiate the framework for KBP in Section 3.5.

### 3.4.1 Problem statement

Let $\mathcal{X}$ be the universe of (relation) instances, $\mathcal{Y} \subseteq \mathcal{X}$ be the unknown subset of correct instances, $X_1, \ldots X_m \subseteq \mathcal{X}$ be the predictions for $m$ systems, and let $Y_i = X_i \cap \mathcal{Y}$. Let $X = \bigcup_{i=1}^m X_i$ and $Y = \bigcup_{i=1}^m Y_i$. Let $f(x) \stackrel{\text{def}}{=} \mathbb{I}[x \in \mathcal{Y}]$ and $g_i(x) = \mathbb{I}[x \in X_i]$, then the precision, $\pi_i$, and recall, $r_i$, of the set of predictions $X_i$ is

$$\pi_i \stackrel{\text{def}}{=} \mathbb{E}_{x \sim p_i}[f(x)] \qquad\qquad r_i \stackrel{\text{def}}{=} \mathbb{E}_{x \sim p_0}[g_i(x)],$$

where $p_i$ is a distribution over $X_i$ and $p_0$ is a distribution over $\mathcal{Y}$. We assume that $p_i$ is known, e.g. the uniform distribution over $X_i$ and that we know $p_0$ up to normalization

constant and can sample from it.

In on-demand evaluation, we can query $f(x)$ (e.g. labeling an instance) or draw a sample from $p_0$; typically, querying $f(x)$ is significantly cheaper than sampling from $p_0$. We obtain prediction sets $X_1, \ldots, X_m$ sequentially as the systems are submitted for evaluation. Our goal is to estimate $\pi_i$ and $r_i$ for each system $i = 1, \ldots, m$.

### 3.4.2 Simple estimators

We can estimate each $\pi_i$ and $r_i$ independently with simple Monte Carlo integration. Let $\hat{X}_1, \ldots, \hat{X}_m$ be multi-sets of $n_1, \ldots, n_j$ i.i.d. samples from $X_1, \ldots, X_m$ respectively, and let $\hat{Y}_0$ be a multi-set of $n_0$ samples drawn from $\mathcal{Y}$. Then, the simple estimators for precision and recall are:

$$\hat{\pi}_i^{(\text{simple})} = \frac{1}{n_i} \sum_{x \in \hat{X}_i} f(x) \qquad\qquad \hat{r}_i^{(\text{simple})} = \frac{1}{n_0} \sum_{x \in \hat{Y}_0} g_i(x).$$

### 3.4.3 Joint estimators

The simple estimators are unbiased but have wastefully large variance because evaluating a new system does not leverage labels acquired for previous systems.

On-demand evaluation with the joint estimator works as follows: First $\hat{Y}_0$ is randomly sampled from $\mathcal{Y}$ once when the evaluation framework is launched. For every new set of predictions $X_m$ submitted for evaluation, the minimum number of samples $n_m$ required to accurately evaluate $X_m$ is calculated based on the current evaluation data, $\hat{Y}_0$ and $\hat{X}_1, \ldots, \hat{X}_{m-1}$. Then, the set $\hat{X}_m$ is added to the evaluation data by evaluating $f(x)$ on $n_m$ samples drawn from $X_m$. Finally, estimates $\pi_i$ and $r_i$ are updated for each system $i = 1, \ldots, m$ using the joint estimators that will be defined next. In the rest of this section, we will answer the following three questions:

1. How can we use all the samples $\hat{X}_1, \ldots \hat{X}_m$ when estimating the precision $\pi_i$ of system $i$?

2. How can we use all the samples $\hat{X}_1, \ldots, \hat{X}_m$ with $\hat{Y}_0$ when estimating recall $r_i$?

3. Finally, to form $\hat{X}_m$, how many samples should we draw from $X_m$ given existing samples and $\hat{X}_1, \ldots, \hat{X}_{m-1}$ and $\hat{Y}_0$?

**Estimating precision jointly.** Intuitively, if two systems have very similar predictions $X_i$ and $X_j$, we should be able to use samples from one to estimate precision on the other. However, it might also be the case that $X_i$ and $X_j$ only overlap on a small region, in which case the samples from $X_j$ do not accurately represent instances in $X_i$ and could lead to a biased estimate. We address this problem by using importance sampling (Owen, 2013), a standard statistical technique for estimating properties of one distribution using samples from another distribution.

In importance sampling, if $\hat{X}_i$ is sampled from $q_i$, then $\frac{1}{n_i} \sum_{x \in \hat{X}_i} \frac{p_i(x)}{q_i(x)} f(x)$ is an unbiased estimate of $\pi_i$. We would like the proposal distribution $q_i$ to both leverage samples from all $m$ systems and be tailored towards system $i$. To this end, we first define a distribution over systems $j$, represented by probabilities $w_{ij}$. Then, define $q_i$ as sampling a $j$ and drawing $x \sim p_j$; formally $q_i(x) = \sum_{j=1}^{m} w_{ij} p_j(x)$.

We note that $q_i(x)$ not only significantly differs between systems, but also changes as new systems are added to the evaluation pool. Unfortunately, the standard importance sampling procedure requires us to draw and use samples from each distribution $q_i(x)$ independently and thus can not effectively reuse samples drawn from different distributions. To this end, we introduce a practical refinement to the importance sampling procedure: we independently draw $n_j$ samples according to $p_j(x)$ from each of the $m$ systems independently and then numerically integrate over these samples using the weights $w_{ij}$ to "mix" them appropriately to produce and unbiased estimate of $\pi_i$ while reducing variance. Formally, we define the *joint precision estimator*:

$$\hat{\pi}_i^{(\text{joint})} \overset{\text{def}}{=} \sum_{j=1}^{m} \frac{w_{ij}}{n_j} \sum_{x \in \hat{X}_j} \frac{p_i(x) f(x)}{q_i(x)},$$

where each $\hat{X}_j$ consists of $n_j$ i.i.d. samples drawn from $p_j$.

It is a hard problem to determine what the optimal mixing weights $w_{ij}$ should be. However, we can formally verify that if $X_i$ and $X_j$ are disjoint, then $w_{ij} = 0$ minimizes the variance of $\pi_i$, and if $X_i = X_j$, then $w_{ij} \propto n_j$ is optimal. This motivates the following heuristic choice which interpolates between these two extremes: $w_{ij} \propto n_j \sum_{x \in \mathcal{X}} p_j(x)p_i(x)$.

**Estimating recall jointly.** The recall of system $i$ can be expressed can be expressed as a product $r_i = \theta \nu_i$, where $\theta$ is the *recall of the pool*, which measures the fraction of all positive instances predicted by the pool (any system), and $\nu_i$ is the *pooled recall of system $i$*, which measures the fraction of the pool's positive instances predicted by system $i$. Letting $g(x) \overset{\text{def}}{=} \mathbb{I}[x \in X]$, we can define these as:

$$\nu_i \overset{\text{def}}{=} \mathbb{E}_{x \sim p_0}[g_i(x) \mid x \in X] \qquad\qquad \theta \overset{\text{def}}{=} \mathbb{E}_{x \sim p_0}[g(x)].$$

We can estimate $\theta$ analogous to the simple recall estimator $\hat{r}_i$, except we use the pool $g$ instead a system $g_i$. For $\nu_i$, the key is to leverage the work from estimating precision. We already evaluated $f(x)$ on $\hat{X}_i$, so we can compute $\hat{Y}_i \overset{\text{def}}{=} \hat{X}_i \cap \mathcal{Y}$ and form the subset $\hat{Y} = \bigcup_{i=1}^m \hat{Y}_i$. $\hat{Y}$ is an approximation of $\mathcal{Y}$ whose bias we can correct through importance reweighting. We then define estimators as follows:

$$\hat{\nu}_i \overset{\text{def}}{=} \frac{\sum_{j=1}^m \frac{w_{ij}}{n_j} \sum_{x \in \hat{Y}_j} \frac{p_0(x)g_i(x)}{q_i(x)}}{\sum_{j=1}^m \frac{w_{ij}}{n_j} \sum_{x \in \hat{Y}_j} \frac{p_0(x)}{q_i(x)}}$$

$$\hat{r}_i^{(\text{joint})} \overset{\text{def}}{=} \hat{\theta}\hat{\nu}_i \quad \hat{\theta} \overset{\text{def}}{=} \frac{1}{n_0} \sum_{x \in \hat{Y}_0} g(x).$$

where $q_i$ and $w_{ij}$ are the same as before.

**Adaptively choosing the number of samples.** Finally, a desired property for on-demand evaluation is to label new instances only when the current evaluation data is insufficient, e.g. when a new set of predictions $X_m$ contains many instances not covered by other systems. We can measure how well the current evaluation set covers the predictions $X_m$ by using a conservative estimate of the variance of $\hat{\pi}_m^{(\text{joint})}$.[6] In particular, the variance of $\hat{\pi}_m^{(\text{joint})}$ is

---

[6]Further details can be found in Appendix **??** of the supplementary material.

a monotonically decreasing function in $n_m$, the number of samples drawn from $X_m$. We can easily solve for the minimum number of samples required to estimate $\hat{\pi}_m^{(\text{joint})}$ within a confidence interval $\epsilon$ by using the bisection method (Burden and Faires, 1985).

## 3.5 On-demand evaluation for KBP

Applying the on-demand evaluation framework to a task requires us to answer three questions:

1. What is the desired distribution over system predictions $p_i$?

2. How do we label an instance $x$, i.e. check if $x \in \mathcal{Y}$?

3. How do we sample from the unknown set of true instances $x \sim p_0$?

In this section, we present practical implementations for knowledge base population.

### 3.5.1 Sampling from system predictions

Both the official TAC-KBP evaluation and the on-demand evaluation we propose use micro-averaged precision and recall as metrics. However, in the official evaluation, these metrics are computed over a fixed set of evaluation entities chosen by LDC annotators, resulting in two problems: (a) defining evaluation entities requires human intervention and (b) typically a large source of variability in evaluation scores comes from not having enough evaluation entities (see e.g. (Webber, 2010)). In our methodology, we replace manually chosen evaluation entities by sampling entities from each system's output according $p_i$. In effect, $p_i$ makes explicit the decision process of the annotator who chooses evaluation entities.

Identifying a reasonable distribution $p_i$ is an important implementation decision that depends on what one wishes to evaluate. Our goal for the on-demand evaluation service we have implemented is to ensure that KBP systems are fairly evaluated on diverse subjects and predicates, while at the same time, ensuring that entities with multiple relations are represented to measure completeness of knowledge base entries. As a result, we propose a distribution that is inversely proportional to the frequency of the subject and predicate

and is proportional to the number of unique relations identified for an entity (to measure knowledge base completeness).

### 3.5.2 Labeling predicted instances

We label predicted relation instances by presenting the instance's provenance to crowd-workers and asking them to identify if a relation holds between the identified subject and object mentions (Figure 3.4a). Crowdworkers are also asked to link the subject and object mentions to their canonical mentions within the document and to pages on Wikipedia, if possible, for entity linking. On average, we find that crowdworkers are able to perform this task in about 20 seconds, corresponding to about $0.05 per instance. We requested 5 crowd-workers to annotate a small set of 200 relation instances from the 2015 TAC-KBP corpus and measured a substantial inter-annotator agreement with a Fleiss' kappa of 0.61 with 3 crowdworkers and 0.62 with 5. Consequently, we take a majority vote over 3 workers in subsequent experiments.

### 3.5.3 Sampling true instances

Sampling from the set of true instances $\mathcal{Y}$ is difficult because we can't even enumerate the elements of $\mathcal{Y}$. As a proxy, we assume that relations are identically distributed across documents and have crowdworkers annotate a random subset of documents for relations using an interface we developed (Figure 3.4b). Crowdworkers begin by identifying every mention span in a document. For each mention, they are asked to identify its type, canonical mention within the document and associated Wikipedia page if possible. They are then presented with a separate interface to label predicates between pairs of mentions within a sentence that were identified earlier.

We compare crowdsourced annotations against those of expert annotators using data from the TAC KBP 2015 EDL task on 10 randomly chosen documents. We find that 3 crowdworkers together identify 92% of the entity spans identified by expert annotators, while 7 crowdworkers together identify 96%. When using a token-level majority vote to identify entities, 3 crowdworkers identify about 78% of the entity spans; this number does

not change significantly with additional crowdworkers. We also measure substantial token-level inter-annotator agreement using Fleiss' kappa for identifying typed mention spans ($\kappa = 0.83$), canonical mentions ($\kappa = 0.75$) and entity links ($\kappa = 0.75$) with just three workers. Based on this analysis, we use token-level majority over 3 workers in subsequent experiments.

The entity annotation interface is far more involved and takes on average about 13 minutes per document, corresponding to about $2.60 per document, while the relation annotation interface takes on average about $2.25 per document. Because documents vary significantly in length and complexity, we set rewards for each document based on the number of tokens (.75c per token) and mention pairs (5c per pair) respectively. With 3 workers per document, we paid about $15 per document on average. Each document contained an average 9.2 relations, resulting in a cost of about $1.61 per relation instance. We note that this is about ten times as much as labeling a relation instance.

## 3.6 Evaluation

Let us now see how well on-demand evaluation works in practice. We begin by empirically studying the bias and variance of the joint estimator proposed in Section 4.3 and find it is able to correct for pooling bias while significantly reducing variance in comparison with the simple estimator. We then demonstrate that on-demand evaluation can serve as a practical replacement for the TAC KBP evalueations by piloting a new evaluation service we have developed to evaluate three distinct systems on TAC KBP 2016 document corpus.

### 3.6.1 Bias and variance of the on-demand evaluation.

Once again, we use the labeled system predictions from the TAC KBP 2015 evaluation and treat them as an exhaustively annotated dataset. To evaluate the pooling methodology we construct an evaluation dataset using instances found by human annotators and labeled instances pooled from 9 randomly chosen teams (i.e. half the total number of participating teams), and use this dataset to evaluate the remaining 9 teams. On average, the pooled evaluation dataset contains between 5,000 and 6,000 labeled instances and evaluates 34

different systems (since each team may have submitted multiple systems). Next, we evaluated sets of 9 randomly chosen teams with our proposed simple and joint estimators using a total of 5,000 samples: about 150 of these samples are drawn from $\mathcal{Y}$, i.e. the full TAC KBP 2015 evaluation data, and 150 samples from each of the systems being evaluated.

We repeat the above simulated experiment 500 times and compare the estimated precision and recall with their true values (Figure 3.4). The simulations once again highlights that the pooled methodology is biased, while the simple and joint estimators are not. Furthermore, the joint estimators significantly reduce variance relative to the simple estimators: the median 90% confidence intervals reduce from 0.14 to 0.06 precision and from 0.14 to 0.08 for recall.

### 3.6.2   Number of samples required by on-demand evaluation

Separately, we evaluate the efficacy of the adaptive sample selection method described in Section 3.4.3 through another simulated experiment. In each trial of this experiment, we evaluate the top 40 systems in random order. As each subsequent system is evaluated, the number of samples to pick from the system is chosen to meet a target variance and added to the current pool of labeled instances. To make the experiment more interpretable, we choose the target variance to correspond with the estimated variance of having 500 samples. Figure 3.4 plots the results of the experiment. The number of samples required to estimate systems quickly drops off from the benchmark of 500 samples as the pool of labeled instances covers more systems. This experiment shows that on-demand evaluation using joint estimation can scale up to an order of magnitude more submissions than a simple estimator for the same cost.

### 3.6.3   A mock evaluation for TAC KBP 2016

We have implemented the on-demand evaluation framework described here as an evaluation service to which researchers can submit their own system predictions. As a pilot of the service, we evaluated three relation extraction systems that also participated in the official 2016 TAC KBP competition. Each system uses Stanford CoreNLP (Manning et al., 2014) to identify entities, the Illinois Wikifier (Ratinov et al., 2011) to perform entity linking and

a combination of a rule-based system (P), a logistic classifier (L), and a neural network classifier (N) for relation extraction. We used 15,000 Newswire documents from the 2016 TAC KBP evaluation as our document corpus. In total, 100 documents were exhaustively annotated for about \$2,000 and 500 instances from each system were labeled for about \$150 each. Evaluating all three system only took about 2 hours.

Figure 3.4f reports scores obtained through on-demand evaluation of these systems as well as their corresponding official TAC evaluation scores. While the relative ordering of systems between the two evaluations is the same, we note that precision and recall as measured through on-demand evaluation are respectively higher and lower than the official scores. This is to be expected because on-demand evaluation measures precision using each systems output as opposed to an externally defined set of evaluation entities. Likewise, recall is measured using exhaustive annotations of relations within the corpus instead of annotations from pooled output in the official evaluation.

## 3.7 Related work

The subject of pooling bias has been extensively studied in the information retrieval (IR) community starting with Zobel (1998), which examined the effects of pooling bias on the TREC AdHoc task, but concluded that pooling bias was not a significant problem. However, when the topic was later revisited, Buckley et al. (2007) identified that the reason for the small bias was because the submissions to the task were too similar; upon repeating the experiment using a novel system as part of the TREC Robust track, they identified a 23% point drop in average precision scores![7]

Many solutions to the pooling bias problem have been proposed in the context of information retrieval, e.g. adaptively constructing the pool to collect relevant data more cost-effectively (Zobel, 1998; Cormack et al., 1998; Aslam et al., 2006), or modifying the scoring metrics to be less sensitive to unassessed data (Buckley and Voorhees, 2004; Sakai and Kando, 2008; Aslam et al., 2006). Many of these ideas exploit the ranking of documents in IR which does not apply to KBP. While both Aslam et al. (2006) and Yilmaz et al. (2008) estimate evaluation metrics by using importance sampling estimators, the techniques they

---

[7]For the interested reader, Webber (2010) presents an excellent survey of the literature on pooling bias.

propose require knowing the set of all submissions beforehand. In contrast, our on-demand methodology can produce unbiased evaluation scores for new development systems as well.

There have been several approaches taken to crowdsource data pertinent to knowledge base population (Vannella et al., 2014; Angeli et al., 2014; He et al., 2015; Liu et al., 2016a). The most extensive annotation effort is probably Pavlick et al. (2016), which crowdsources a knowledge base for gun-violence related events. In contrast to previous work, our focus is on *evaluating systems*, not collecting a dataset. Furthermore, our main contribution is not a large dataset, but an evaluation service that allows anyone to use crowdsourcing predictions made by their system.

## 3.8 Discussion

Over the last ten years of the TAC KBP task, the gap between human and system performance has barely narrowed despite the community's best efforts: top automated systems score less than 36% $F_1$ while human annotators score more than 60%. In this paper, we've shown that the current evaluation methodology may be a contributing factor because of its bias against novel system improvements. The new on-demand framework proposed in this work addresses this problem by obtaining human assessments of new system output through crowdsourcing. The framework is made economically feasible by carefully sampling output to be assessed and correcting for sample bias through importance sampling.

Of course, simply providing better evaluation scores is only part of the solution and it is clear that better datasets are also necessary. However, the very same difficulties in scale that make evaluating KBP difficult also make it hard to collect a high quality dataset for the task. As a result, existing datasets (Angeli et al., 2014; Adel et al., 2016) have relied on the output of existing systems, making it likely that they exhibit the same biases against novel systems that we've discussed in this paper. We believe that providing a fair and standardized evaluation platform as a service allows researchers to exploit such datasets and while still being able to accurately measure their performance on the knowledge base population task.

There are many other tasks in NLP that are even harder to evaluate than KBP. Existing evaluation metrics for tasks with a generation component—such as summarization or

dialogue—leave much to be desired. We believe that adapting the ideas of this paper to those tasks is a fruitful direction, as progress of a research community is strongly tied to the fidelity of evaluation.

| | Median bias | | |
|---|---|---|---|
| | Precision | Recall | Macro $F_1$ |
| Official | 17.93% | 17.00% | 15.51% |
| `anydoc` | 2.34% | 1.93% | 2.05% |

Figure 3.3: Median pooling bias (difference between pooled and unpooled scores) on the top 40 systems of TAC KBP 2015 evaluation using the official and `anydoc` scores. The bias is much smaller for the lenient `anydoc` metric, but even so, it is larger than the largest difference between adjacent systems (1.5% $F_1$) and typical system improvements (around 1% $F_1$).

(a)

(b)

(c)

(d)

(e)

| Sys. | $P$ | $R$ | $F_1$ |
|------|-----|-----|-------|
| *TAC KBP evaluation* | | | |
| P | 47.6% | 11.0% | 17.9% |
| P+L | 35.5% | 18.4% | 24.2% |
| P+L+N | 26.3% | 27.0% | 26.6% |
| *On-demand evaluation* | | | |
| P | 74.7% | 5.8% | 10.8% |
| P+L | 54.7% | 7.6% | 13.3% |
| P+L+N | 34.0% | 9.8% | 15.2% |

(f)

Figure 3.4: **(a, b):** Interfaces for annotating relations and entities respectively. **(c, d):** A comparison of bias for the pooling, simple and joint estimators on the TAC KBP 2015 challenge. Each point in the figure is a mean of 500 repeated trials; dotted lines show the 90% quartile. Both the simple and joint estimators are unbiased, and the joint estimator is able to significantly reduce variance. **(e):** A comparison of the number of samples used to estimate scores under the fixed and adaptive sample selection scheme. Each faint line shows the number of samples used during a single trial, while solid lines show the mean over 100 trials. The dashed line shows a square-root relationship between the number of systems evaluated and the number of samples required. Thus joint estimation combined with adaptive sample selection can reduce the number of labeled annotations required by

# Chapter 4

# The price of debiasing automatic metrics in natural language evaluation

## Abstract

For evaluating generation systems, automatic metrics such as BLEU cost nothing to run but have been shown to correlate poorly with human judgment, leading to systematic bias against certain model improvements. On the other hand, averaging human judgments, the unbiased gold standard, is often too expensive. In this paper, we use control variates to combine automatic metrics with human evaluation to obtain an unbiased estimator with lower cost than human evaluation alone. In practice, however, we obtain only a 7–13% cost reduction on evaluating summarization and open-response question answering systems. We then prove that our estimator is optimal: there is no unbiased estimator with lower cost. Our theory further highlights the two fundamental bottlenecks—the automatic metric and the prompt shown to human evaluators—both of which need to be improved to obtain greater cost savings.

## 4.1   Introduction

In recent years, there has been an increasing interest in tasks that require generating natural language, including abstractive summarization (Nallapati et al., 2016), open-response question answering (Nguyen et al., 2016; Kočisky et al., 2017), image captioning (Lin et al., 2014), and open-domain dialogue (Lowe et al., 2017b). Unfortunately, the evaluation of these systems remains a thorny issue because of the diversity of possible correct responses. As the gold standard of performing human evaluation is often too expensive, there has been a large effort developing automatic metrics such as BLEU (Papineni et al., 2002), ROUGE (Lin and Rey, 2004), METEOR (Lavie and Denkowski, 2009; Denkowski and Lavie, 2014) and CiDER (Vedantam et al., 2015). However, these have shown to be biased, correlating poorly with human metrics across different datasets and systems (Liu et al., 2016b; Novikova et al., 2017).



(a) System-level correlation on the MS MARCO task

(b) Instance-level correlation for the `fastqa` system

Figure 4.1:   (a) At a system-level, automatic metrics (ROUGE-L) and human judgment correlate well, but (b) the instance-level correlation plot (where each point is a system prediction) shows that the instance-level correlation is quite low ($\rho = 0.31$). As a consequence, if we try to locally improve systems to produce better answers ($\triangleright$ in (a)), they do not significantly improve ROUGE scores and vice versa ($\triangle$).

Can we combine automatic metrics and human evaluation to obtain an *unbiased* estimate at *lower cost* than human evaluation alone? In this paper, we propose a simple estimator based on control variates (Ripley, 2009), where we average differences between

human judgments and automatic metrics rather than averaging the human judgments alone. Provided the two are correlated, our estimator will have lower variance and thus reduce cost.

We prove that our estimator is *optimal* in the sense that no unbiased estimator using the same automatic metric can have lower variance. We also analyze its data efficiency (equivalently, cost savings)—the factor reduction in number of human judgments needed to obtain the same accuracy versus naive human evaluation—and show that it depends solely on two factors: (a) the annotator variance (which is a function of the human evaluation prompt) and (b) the correlation between human judgments and the automatic metric. This factorization allows us to calculate typical and best-case data efficiencies and accordingly refine the evaluation prompt or automatic metric.

Finally, we evaluate our estimator on state-of-the-art systems from two tasks, summarization on the CNN/Daily Mail dataset Hermann et al. (2015); Nallapati et al. (2016) and open-response question answering on the MS MARCOv1.0 dataset (Nguyen et al., 2016). To study our estimators offline, we preemptively collected 10,000 human judgments which cover several tasks and systems.[1] As predicted by the theory, we find that the data efficiency depends not only on the correlation between the human and automatic metrics, but also on the evaluation prompt. If the automatic metric had perfect correlation, our data efficiency would be around 3, while if we had noiseless human judgments, our data efficiency would be about 1.5. In reality, the reduction in cost we obtained was only about 10%, suggesting that improvements in both automatic metric and evaluation prompt are needed. As one case study in improving the latter, we show that, when compared to a Likert survey, measuring the amount of post-editing needed to fix a generated sentence reduced the annotator variance by three-fold.

## 4.2   Bias in automatic evaluation

It is well understood that current automatic metrics tend to correlate poorly with human judgment at the instance-level. For example, Novikova et al. (2017) report correlations

---

[1] An anonymized version of this data and the annotation interfaces used can be found at https://bit.ly/price-of-debiasing.

| Question and reference answer | System answer (System; `Corr` / ROUGE-L) |
|---|---|
| *Examples where system is correct and ROUGE-L > 0.5 (19.6% or 285 of 1455 unique responses)* | |
| **Q.** what is anti-mullerian hormone<br>**A.** Anti-Mullerian Hormone (AMH) is a protein hormone produced by granulosa cells (cells lining the egg sacs or follicles) within the ovary. | it is a protein hormone produced by granulosa cells (cells lining the egg sacs or follicles) within the ovary. (`snet.ens`; ✓ / 0.86) |
| *Examples where system is incorrect and ROUGE-L > 0.5 (1.3% or 19 of 1455 unique responses)* | |
| **Q.** at what gestational age can you feel a fetus move<br>**A.** 37 to 41 weeks *(incorrect reference answer)* | 37 to 41 weeks (`fastqa, fastqa.ext`; ✗ / 1.0) |
| *Examples where system is correct and ROUGE-L < 0.5 (56.0% or 815 of 1455 unique responses)* | |
| **Q.** what is the definition of onomatopoeia<br>**A.** It is defined as a word, which imitates the natural sounds of a thing. | the naming of a thing or action by a vocal imitation of the sound associated with it (as buzz, hiss). (`fastqa`; ✓ / 0.23) |
| *Examples where system is incorrect and ROUGE-L < 0.5 (23.1% or 336 of 1455 unique responses)* | |
| **Q.** what kind root stem does a dandelion have<br>**A.** Fibrous roots and hollow stem. | vitamin a, vitamin c, vitamin d and vitamin b complex, as well as zinc, iron and potassium. (`snet, snet.ens`; ✗ / 0.09) |

(a) **MS MARCO.** Human annotators rated answer correctness (`AnyCorrect`) and the automatic metric used is ROUGE-L (higher is better).

| Reference summary | System summary (System; `Edit` / VecSim) |
|---|---|
| *Examples where system Edit < 0.3 and VecSim > 0.5 (53.9% or 1078 of 2000 responses)* | |
| Bhullar is set to sign a ■-day contract with the Kings. The ■-year-old will become the NBA's first player of Indian descent. Bhullar will be on the roster when the Kings host New Orleans Pelicans. | ~~Bhullar and~~**The** Kings are signing Bhullar to a ■-day contract. The ■-year-old will be on the roster on friday when David Wear's ■-season contract expires thursday. Bhullar is set to become the NBA's first player of Indian descent. (`ml`; 0.13 / 0.82) |
| *Examples where system Edit > 0.3 and VecSim > 0.5 (18.0% or 360 of 2000 responses)* | |
| The Direct Marketing Commission probing B2C Data and Data Bubble. Investigating whether they breached rules on the sale of private data. Chief commissioner described allegations made about firms as 'serious'. | ■ ~~Data obtained by the Mail's marketing commission said it would probe both companies over claims that they had breached the rules on the sale of private data.~~ The FSA said it would probe both companies over claims they had breached the rules on the sale of private data. (`se2seq`; 1.00 / 0.72) |
| *Examples where system Edit < 0.3 and VecSim < 0.5 (14.5% or 290 of 2000 responses)* | |
| Death toll rises to more than ■. Pemba Tamang, ■, shows no apparent signs of serious injury after rescue. Americans | ~~Six of~~ **Despite** Nepal's tragedy, life triumphed in Kathmandu's hard-hit neighborhoods. Rescuers pulled an 15-year-old from the rubble of a multi- |

less than $0.3$ for a large suite of word-based and grammar-based evaluation methods on a generation task. Similarly, Liu et al. (2016b) find correlations less than $0.35$ for automatic metrics on a dialog generation task in one domain, but find correlations with the same metric dropped significantly to less than $0.16$ when used in another domain. Still, somewhat surprisingly, several automatic metrics have been found to have high *system-level* correlations (Novikova et al., 2017). What, then, are the implications of having a low instance-level correlation?

As a case study, consider the task of open-response question answering: here, a system receives a human-generated question and must *generate* an answer from some given context, e.g. a document or several webpages. We collected the responses of several systems on the MS MARCOv1 dataset (Nguyen et al., 2016) and crowdsourced human evaluations of the system output (see Section 4.4 for details).

The instance-level correlation (Figure 4.1b) is only $\rho = 0.31$. A closer look at the instance-level correlation reveals that while ROUGE is able to correctly assign low scores to bad examples (lower left), it is bad at judging good examples and often assigns them low ROUGE scores (lower right)—see Table 4.1 for examples. This observation agrees with a finding reported in Novikova et al. (2017) that automatic metrics correlate better with human judgments on bad examples than average or good examples.

Thus, as Figure 4.1(a) shows, we can improve low-scoring ROUGE examples without improving their human judgment ($\triangle$) and vice versa ($\triangleright$). Indeed, Conroy and Dang (2008) report that summarization systems were optimized for ROUGE during the DUC challenge (Dang, 2006) until they were indistinguishable from the ROUGE scores of human-generated summaries, but the systems had hardly improved on human evaluation. Hill-climbing on ROUGE can also lead to a system that does worse on human scores, e.g. in machine translation (Wu et al., 2016). Conversely, genuine quality improvements might not be reflected in improvements in ROUGE. This bias also appears in pool-based evaluation for knowledge base population (Chaganty et al., 2017). Thus the problems with automatic metrics clearly motivate the need for human evaluation, but can we still use the automatic metrics somehow to save costs?

## 4.3 Statistical estimation for unbiased evaluation

We will now formalize the problem of combining human evaluation with an automatic metric. Let $\mathcal{X}$ be a set of inputs (e.g., articles), and let $S$ be the *system* (e.g. for summarization), which takes $x \in \mathcal{X}$ and returns output $S(x)$ (e.g. a summary). Let $\mathcal{Z} = \{(x, S(x)) : x \in \mathcal{X}\}$ be the set of system predictions. Let $Y(z)$ be the random variable representing the human judgment according to some evaluation prompt (e.g. grammaticality or correctness), and define $f(z) = \mathbb{E}[Y(z)]$ to be the (unknown) *human metric* corresponding to averaging over an infinite number of human judgments. Our goal is to estimate the average across all examples:

$$\mu \stackrel{\text{def}}{=} \mathbb{E}_z[f(z)] = \frac{1}{|\mathcal{Z}|} \sum_{z \in \mathcal{Z}} f(z) \tag{4.1}$$

with as few queries to $Y$ as possible.

Let $g$ be an automatic metric (e.g. ROUGE), which maps $z$ to a real number. We assume evaluating $g(z)$ is free. The central question is how to use $g$ in conjunction with calls to $Y$ to produce an unbiased estimate $\hat{\mu}$ (that is, $\mathbb{E}[\hat{\mu}] = \mu$). In this section, we will construct a simple estimator based on control variates (Ripley, 2009), and prove that it is minimax optimal.

### 4.3.1 Sample mean

We warm up with the most basic unbiased estimate, the sample mean. We sample $z^{(1)}, \ldots, z^{(n)}$ independently with replacement from $\mathcal{Z}$. Then, we sample each human judgment $y^{(i)} = Y(z^{(i)})$ independently.[2] Define the estimator to be $\hat{\mu}_{\text{mean}} = \frac{1}{n} \sum_{i=1}^{n} y^{(i)}$. Note that $\hat{\mu}_{\text{mean}}$ is unbiased ($\mathbb{E}[\hat{\mu}_{\text{mean}}] = \mu$).

We can define $\sigma_f^2 \stackrel{\text{def}}{=} \text{Var}(f(z))$ as the variance of the human metric and $\sigma_a^2 \stackrel{\text{def}}{=} \mathbb{E}_z[\text{Var}(Y(z))]$

---

[2]Note that this independence assumption isn't quite true in practice since we do not control who annotates our data.

as the variance of human judgment averaged over $\mathcal{Z}$. By the law of total variance, the variance of our estimator is

$$\mathrm{Var}(\hat{\mu}_{\mathrm{mean}}) = \frac{1}{n}(\sigma_f^2 + \sigma_a^2). \tag{4.2}$$

## 4.3.2 Control variates estimator

Now let us see how an automatic metric $g$ can reduce variance. If there is no annotator variance ($\sigma_a^2 = 0$) so that $Y(z) = f(z)$, we should expect the variance of $f(z) - g(z)$ to be lower than the variance of $f(z)$, assuming $g$ is correlated with $f$—see Figure 4.2 for an illustration.

The actual control variates estimator needs to handle noisy $Y(z)$ (i.e. $\sigma_a^2 > 0$) and guard against a $g(z)$ with low correlation. Let us standardize $g$ to have zero mean and unit variance, because we have assumed it is free to evaluate. As before, let $z^{(1)}, \ldots, z^{(n)}$ be independent samples from $\mathcal{Z}$ and draw $y^{(i)} = Y(z^{(i)})$ independently as well. We define the *control variates estimator* as

$$\hat{\mu}_{\mathrm{cv}} = \frac{1}{n} \sum_{i=1}^{n} y^{(i)} - \alpha g(z^{(i)}), \tag{4.3}$$

where

$$\alpha \stackrel{\mathrm{def}}{=} \mathrm{Cov}(f(z), g(z)). \tag{4.4}$$

Intuitively, we have averaged over $y^{(i)}$ to handle the noise introduced by $Y(z)$, and scaled $g(z)$ to prevent an uncorrelated automatic metric from introducing too much noise.

An important quantity governing the quality of an automatic metric $g$ is the correlation between $f(z)$ and $g(z)$ (recall that $g$ has unit variance):

$$\rho \stackrel{\mathrm{def}}{=} \frac{\alpha}{\sigma_f}. \tag{4.5}$$

We can show that among all distributions with fixed $\sigma_f^2$, $\sigma_a^2$, and $\alpha$ (equivalently $\rho$), this estimator is minimax optimal, i.e. it has the least variance among all unbiased estimators:

Figure 4.2: The samples from $f(z)$ have a higher variance than the samples from $f(z) - g(z)$ but the same mean. This is the key idea behind using control variates to reduce variance.

**Theorem 1.** *Among all unbiased estimators that are functions of $y^{(i)}$ and $g(z^{(i)})$, and for all distributions with a given $\sigma_f^2$, $\sigma_a^2$, and $\alpha$,*

$$\text{Var}(\hat{\mu}_{cv}) = \frac{1}{n}(\sigma_f^2(1 - \rho^2) + \sigma_a^2), \tag{4.6}$$

*and no other estimator has a lower worst-case variance.*

Comparing the variances of the two estimators ((4.2) and (4.6)), we define the *data efficiency* as the ratio of the variances:

$$\text{DE} \overset{\text{def}}{=} \frac{\text{Var}(\hat{\mu}_{\text{mean}})}{\text{Var}(\hat{\mu}_{\text{cv}})} = \frac{1 + \gamma}{1 - \rho^2 + \gamma}, \tag{4.7}$$

where $\gamma \overset{\text{def}}{=} \sigma_a^2/\sigma_f^2$ is the normalized annotator variance. Data efficiency is the key quantity in this paper: it is the multiplicative reduction in the number of samples required when using the control variates estimator $\hat{\mu}_{\text{cv}}$ versus the sample mean $\hat{\mu}_{\text{mean}}$. Figure 4.3 shows the inverse data efficiency contours as a function of the correlation $\rho$ and $\gamma$.

When there is no correlation between human and automatic metrics ($\rho = 0$), the data efficiency is naturally $1$ (no gain). In order to achieve a data efficiency of $2$ (half the labeling cost), we need $|\rho| \geq \sqrt{2}/2 \approx 0.707$. Interestingly, even for an automatic metric with perfect correlation ($\rho = 1$), the data efficiency is still capped by $\frac{1+\gamma}{\gamma}$: unless $\gamma \to 0$ the data efficiency cannot increase unboundedly. Intuitively, even if we knew that $\rho = 1$, $f(z)$ would be undetermined up to a constant additive shift and just estimating the shift would incur a variance of $\frac{1}{n}\sigma_a^2$.

### 4.3.3 Using the control variates estimator

The control variates estimator can be easily integrated into an existing evaluation: we run human evaluation on a random sample of system outputs, automatic evaluation on all the system outputs, and plug in these results into Algorithm 1.

It is vital that we are able to evaluate the automatic metric on a significantly larger set of examples than those with human evaluations to reliably normalize $g(z)$: without these additional examples, it be can shown that the optimal minimax estimator for $\mu$ is simply

Figure 4.3: Inverse data efficiency for various values of $\gamma$ and $\rho$. We need both low $\gamma$ and high $\rho$ to obtain significant gains.

the naive estimate $\hat{\mu}_{\text{mean}}$. Intuitively, this is because estimating the mean of $g(z)$ incurs an equally large variance as estimating $\mu$. In other words, $g(z)$ is only useful if we have additional information about $g$ beyond the samples $\{z^{(i)}\}$.

Algorithm 1 shows the estimator. In practice, we do not know $\alpha = \text{Cov}(f(z), g(z))$, so we use a plug-in estimate $\hat{\alpha}$ in line 3 to compute the estimate $\widetilde{\mu}$ in line 4. We note that estimating $\alpha$ from data does introduce a $O(1/n)$ bias, but when compared to the standard deviation which decays as $\Theta(1/\sqrt{n})$, this bias quickly goes to $0$.

**Proposition 1.** *The estimator $\widetilde{\mu}$ in Algorithm 1 has $O(1/n)$ bias.*

---

**Algorithm 1** Control variates estimator

---

1: **Input:** $n$ human evaluations $y^{(i)}$ on system outputs $z^{(i)}$, *normalized* automatic metric $g$
2: $\overline{y} = \frac{1}{n} \sum_i y^{(i)}$
3: $\hat{\alpha} = \frac{1}{n} \sum_i (y^{(i)} - \overline{y}) g(z^{(i)})$
4: $\widetilde{\mu} = \frac{1}{n} \sum_i y^{(i)} - \hat{\alpha} g(z^{(i)})$
5: **return** $\widetilde{\mu}$

---

An additional question that arises when applying Algorithm 1 is figuring out how many samples $n$ to use. Given a target variance, the number of samples can be estimated using (4.6) with conservative estimates of $\sigma_f^2$, $\sigma_a^2$ and $\rho$. Alternatively, our estimator can be combined with a dynamic stopping rule (Mnih et al., 2008) to stop data collection once we reach a target confidence interval.

### 4.3.4 Discussion of assumptions

We will soon see that empirical instantiations of $\gamma$ and $\rho$ lead to rather underwhelming data efficiencies in practice. In light of our optimality result, does this mean there is no hope for gains? Let us probe our assumptions. We assumed that the human judgments are uncorrelated across different system outputs; it is possible that a more accurate model of human annotators (e.g. Passonneau and Carpenter (2014)) could offer improvements. Perhaps with additional information about $g(z)$ such as calibrated confidence estimates, we

(a) Interface to evaluate language quality on CNN/Daily Mail



(b) Interface to judge answer correctness on MS MARCO
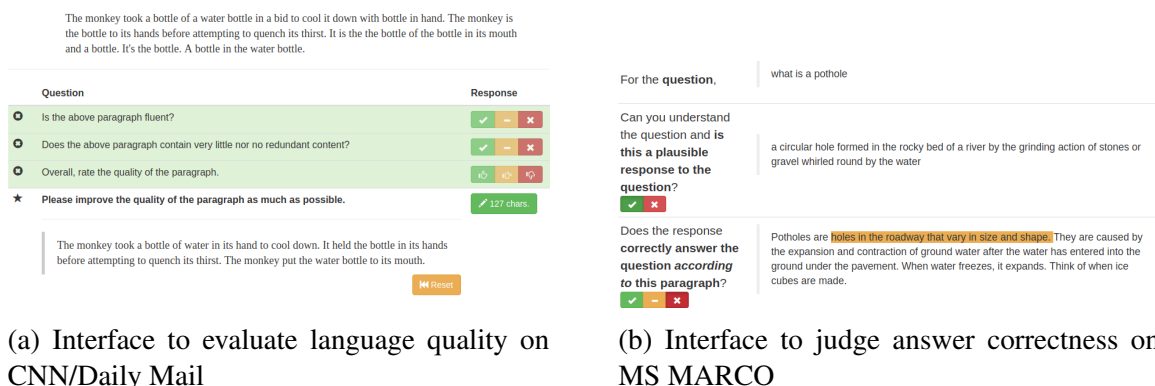
Figure 4.4: Screenshots of the annotation interfaces we used to measure (a) summary language quality on CNN/Daily Mail and (b) answer correctness on MS MARCO tasks.

would be able to sample more adaptively. Of course the most direct routes to improvement involve increasing the correlation of $g$ with human judgments and reducing annotator variance, which we will discuss more later.

## 4.4 Tasks and datasets

In order to compare different approaches to evaluating systems, we first collected human judgments for the output of several automatic summarization and open-response question answering systems using Amazon Mechanical Turk. Details of instructions provided and quality assurance steps taken are provided in Appendix **??** of the supplementary material. In this section, we'll briefly describe how we collected this data.

**Evaluating language quality in automatic summarization.** In automatic summarization, systems must generate a short (on average two or three sentence) summary of an article: for our study, we chose articles from the CNN/Daily Mail (CDM) dataset (Hermann et al., 2015; Nallapati et al., 2016) which come paired with reference summaries in the form of story highlights. We focus on the *language quality* of summaries and leave evaluating content selection to future work.

For each summary, we collected human judgments on a scale from 1–3 (Figure 4.4a) for fluency, (lack of) redundancy, and overall quality of the summary using guidelines from the

| Task | Eval. | $\sigma_a^2$ | $\sigma_f^2$ | $\gamma = \frac{\sigma_a^2}{\sigma_f^2}$ |
|---|---|---|---|---|
| CDM | Fluency | 0.32 | 0.26 | 1.23 |
| CDM | Redund. | 0.26 | 0.43 | 0.61 |
| CDM | Overall | 0.28 | 0.28 | 1.00 |
| **CDM** | **Edit** | **0.07** | **0.18** | **0.36** |
| MS MARCO | AnyCorr. | 0.14 | 0.15 | 0.95 |
| MS MARCO | AvgCorr. | 0.12 | 0.13 | 0.91 |

Table 4.2: A summary of the key statistics, human metric variance ($\sigma_f^2$) and annotator variance ($\sigma_a^2$) for different datasets, CNN/Daily Mail (CDM) and MS MARCO in our evaluation benchmark. We observe that the relative variance ($\gamma$) is fairly high for most evaluation prompts, upper bounding the data efficiency on these tasks. A notable exception is the `Edit` prompt wherein systems are compared on the number of post-edits required to improve their quality.

DUC summarization challenge (Dang, 2006). As an alternate human metric, we also asked workers to post-edit the system's summary to improve its quality, similar to the post-editing step in MT evaluations (Snover et al., 2006). Obtaining judgments costs about $0.15 per summary and this cost rises to about $0.40 per summary for post-editing.

We collected judgments on the summaries generated by the `seq2seq` and `pointer` models of See et al. (2017), the `ml` and `ml+rl` models of Paulus et al. (2018), and the reference summaries.[3] Before presenting the summaries to human annotators, we performed some minimal post-processing: we true-cased and de-tokenized the output of `seq2seq` and `pointer` using Stanford CoreNLP (Manning et al., 2014) and replaced "unknown" tokens in each system with a special symbol (■).

**Evaluating answer correctness.** Next, we look at evaluating the correctness of system outputs in question answering using the MS MARCO question answering dataset (Nguyen et al., 2016). Here, each system is provided with a question and up to 10 paragraphs of context. The system generates open-response answers that do not need to be tied to a span in any paragraph.

We first ask annotators to judge if the output is even plausible for the question, and if

---

[3]All system output was obtained from the original authors through private communication.

yes, ask them identify if it is correct according to each context paragraph. We found that requiring annotators to highlight regions in the text that support their decision substantially improved the quality of the output without increasing costs. Annotations cost $0.40 per system response.[4]

While our goal is to evaluate the correctness of the provided answer, we found that there are often answers which may be correct or incorrect depending on the context. For example, the question "what is a pothole" is typically understood to refer to a hole in a roadway, but also refers to a geological feature (Figure 4.4b). This is reflected when annotators mark one context paragraph to support the given answer but mark another to contradict it. We evaluated systems based on both the average correctness (AvgCorrect) of their answers across all paragraphs as well as whether their answer is correct according to any paragraph (AnyCorrect).

We collected annotations on the systems generated by the `fastqa` and `fastqa_ext` from Weissenborn et al. (2017) and the `snet` and `snet.ens`(emble) models from Tan et al. (2018), along with reference answers. The answers generated by the systems were used without any post-processing. Surprisingly, we found that the correctness of the reference answers (according to the AnyCorrect metric) was only 73.5%, only 2% above that of the leading system (`snet.ens`). We manually inspected 30 reference answers which were annotated incorrectly and found that of those, about 95% were indeed incorrect. However, 62% are actually answerable from some paragraph, indicating that the real ceiling performance on this dataset is around 90% and that there is still room for improvement on this task.

## 4.5   Experimental results

We are now ready to evaluate the performance of our control variates estimator proposed in Section 4.3 using the datasets presented in Section 4.4. Recall that our primary quantity of interest is *data efficiency*, the ratio of the number of human judgments required to estimate the overall human evaluation score for the control variates estimator versus the sample

---

[4]This cost could be significantly reduced if systems also specify which passage they used to generate the answer.

(a) MS MARCO with the `AnyCorrect` prompt



(b) CNN/Daily Mail with the `Edit` prompt

Figure 4.5:   Correlations of different automatic metrics on the MS MARCO and CNN/Daily Mail tasks. Certain systems are more correlated with certain automatic metrics than others, but overall the correlation is low to moderate for most systems and metrics.

mean. We'll briefly review the automatic metrics used in our evaluation before analyzing the results.

**Automatic metrics.**   We consider the following frequently used automatic word-overlap based metrics in our work: **BLEU** (Papineni et al., 2002), **ROUGE** (Lin and Rey, 2004) and **METEOR** (Lavie and Denkowski, 2009). Following Novikova et al. (2017) and Liu et al. (2016b), we also compared a vector-based sentence-similarity using `sent2vec` (Pagliardini et al., 2017) to compare sentences (**VecSim**). Figure 4.5 shows how each of these metrics is correlated with human judgment for the systems being evaluated. Unsurprisingly, the correlation varies considerably across systems, with token-based metrics correlating more strongly for systems that are more extractive in nature (`fastqa` and `fastqa_ext`).

**Results.**   [5]

In Section 4.3 we proved that the control variates estimator is not only unbiased but also has the least variance among other unbiased estimators. Figure 4.6 plots the width of the 80% confidence interval, estimated using bootstrap, measured as a function of the

---

[5]Extended results for other systems, metrics and prompts can be found at https://bit.ly/price-of-debiasing/.

(a) `seq2seq` on CNN/Daily Mail using the `Overall`

(b) `seq2seq` on CNN/Daily Mail using `Edit`

(c) `fastqa_ext` on MS MARCO using `AnyCorrect`
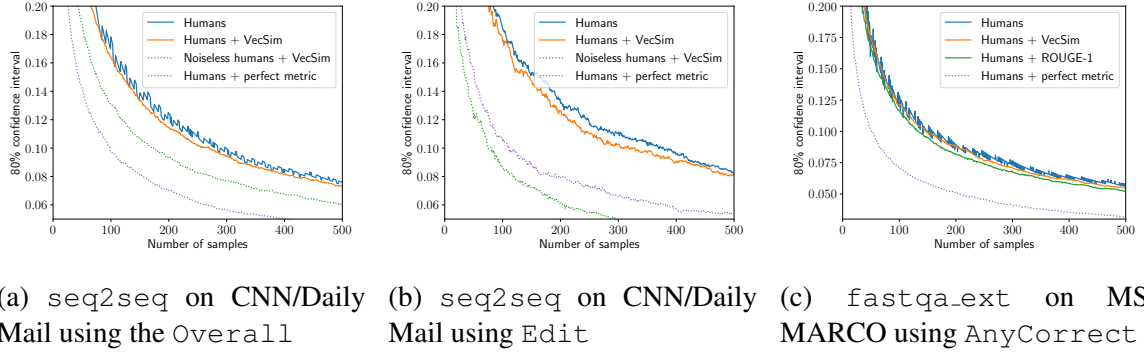
Figure 4.6: 80% bootstrap confidence interval length as a function of the number of human judgments used when evaluating the indicated systems on their respective datasets and prompts. (a) We see a modest reduction in variance (and hence cost) relative to human evaluation by using the VecSim automatic metric with the proposed control variates estimator to estimate `Overall` scores on the CNN/Daily Mail task; the data efficiency (DE) is $1.06$. (b) By improving the evaluation prompt to use `Edits` instead, it is possible to further reduce variance relative to humans (DE is $1.15$). (c) Another way to reduce variance relative to humans is to improve the automatic metric evaluation; here using ROUGE-1 instead of VecSim improves the DE from $1.03$ to $1.16$.

number of samples collected for different tasks and prompts. As expected, the control variates estimator reduces the width of the confidence interval. We measure data efficiency by the averaging of the ratio of squared confidence intervals between the human baseline and control variates estimates. We observe that the data efficiency depends on the task, prompt and system, ranging from about $1.08$ (a 7% cost reduction) to $1.15$ (a 13% cost reduction) using current automatic metrics.

As we showed in Section 4.3, further gains are fundamentally limited by the quality of the evaluation prompts and automatic metrics. Figures 4.6a and 4.6b show how improving the quality of the evaluation prompt from a Likert-scale prompt for quality (`Overall`) to using post-editing (`Edit`) noticeably decreases variance and hence allows better automatic metrics to increase data efficiency. Likewise, Figure 4.6c shows how using a better automatic metric (ROUGE-L instead of VecSim) also reduces variance.

Figure 4.6 also shows the conjectured confidence intervals if we were able to eliminate noise in human judgments (noiseless humans) or have a automatic metric that correlated perfectly with average human judgment (perfect metric). In particular, we use the mean of

all (2–3) humans on each $z$ for the perfect $g(z)$ and use the mean of all humans on each $z$ for the "noiseless" $Y(z)$.

In both cases, we are able to significantly increase data efficiency (i.e. decrease estimator variance). With zero annotator variance and using existing automatic metrics, the data efficiency ranges from 1.42 to 1.69. With automatic metrics with perfect correlation and current variance of human judgments, it ranges from 2.38 to 7.25. Thus, we conclude that it is important not only to improve our automatic metrics but also the evaluation prompts we use during human evaluation.

## 4.6 Related work

In this work, we focus on using existing automatic metrics to decrease the cost of human evaluations. There has been much work on improving the quality of automatic metrics. In particular, there is interest in learning models (Lowe et al., 2017a; Dusek et al., 2017) that are able to optimize for improved correlations with human judgment. However, in our experience, we have found that these learned automatic metrics have trouble generalizing to different systems. The framework we provide allows us to safely incorporate such models into evaluation, exploiting them when their correlation is high but also not introducing bias when it is low.

Our key technical tool is control variates, a standard statistical technique used to reduce the variance of Monte Carlo estimates (Ripley, 2009). The technique has also been used in machine learning and reinforcement learning to lower variance estimates of gradients (Greensmith et al., 2004; Paisley et al., 2012; Ranganath et al., 2014). To the best of our knowledge, we are the first to apply this technique in the context of language evaluation.

Our work also highlights the importance of human evaluation. Chaganty et al. (2017) identified a similar problem of systematic bias in evaluation metrics in the setting of knowledge base population and also propose statistical estimators that relies on human evaluation to correct bias. Unfortunately, their technique relies on having a structured output (relation triples) that are shared between systems and does not apply to evaluating natural language generation. In a similar vein, Chang et al. (2017) dynamically collect human feedback to learn better dialog policies.

## 4.7 Discussion

Prior work has shown that existing automatic metrics have poor instance-level correlation with mean human judgment and that they score many good quality responses poorly. As a result, the evaluation is systematically biased against genuine system improvements that would lead to higher human evaluation scores but not improve automatic metrics. In this paper, we have explored using an automatic metric to decrease the cost of human evaluation without introducing bias. In practice, we find that with current automatic metrics and evaluation prompts data efficiencies are only 1.08–1.15 (7–13% cost reduction). Our theory shows that further improvements are only possible by improving the correlation of the automatic metric and reducing the annotator variance of the evaluation prompt. As an example of how evaluation prompts could be improved, we found that using post-edits of summarizes decreased normalized annotator variance by a factor of three relative to using a Likert scale survey. It should be noted that changing the evaluation prompt also changes the underlying ground truth $f(z)$: it is up to us to find a prompt that still captures the essence of what we want to measure.

Without making stronger assumptions, the control variates estimator we proposed outlines the limitations of unbiased estimation. Where do we go from here? Certainly, we can try to improve the automatic metric (which is potentially as difficult as solving the task) and brainstorming alternative ways of soliciting evaluation (which has been less explored). Alternatively, we could give up on measuring absolute scores, and seek instead to find techniques stably rank methods and thus improve them. As the NLP community tackles increasingly difficult tasks, human evaluation will only become more important. We hope our work provides some clarity on to how to make it more cost effective.

# Chapter 5

# On-the-Job Learning with Bayesian Decision Theory

Our goal is to deploy a high-accuracy system starting with zero training examples. We consider an "on-the-job" setting, where as inputs arrive, we use real-time crowdsourcing to resolve uncertainty where needed and output our prediction when confident. As the model improves over time, the reliance on crowdsourcing queries decreases. We cast our setting as a stochastic game based on Bayesian decision theory, which allows us to balance latency, cost, and accuracy objectives in a principled way. Computing the optimal policy is intractable, so we develop an approx- imation based on Monte Carlo Tree Search. We tested our approach on three datasets—named-entity recognition, sentiment classification, and image classification. On the NER task we obtained more than an order of magnitude reduction in cost compared to full human annotation, while boosting performance relative to the expert provided labels. We also achieve a $8\%$ $F_1$ improvement over having a single human label the whole set, and a $28\%$ $F_1$ improvement over online learning.

'

---

'

Poor is the pupil who does not surpass his master."
– Leonardo da Vinci

## 5.1    Introduction

There are two roads to an accurate AI system today: (i) gather a huge amount of labeled training data (Deng et al., 2009) and do supervised learning (Krizhevsky et al., 2012); or (ii) use crowdsourcing to directly perform the task (Bernstein et al., 2010; Kokkalis et al., 2013). However, both solutions require non-trivial amounts of time and money. In many situations, one wishes to build a new system — e.g., to do Twitter information extraction (Li et al., 2012) to aid in disaster relief efforts or monitor public opinion — but one simply lacks the resources to follow either the pure ML or pure crowdsourcing road.

In this paper, we propose a framework called *on-the-job learning* (formalizing and extending ideas first implemented in (**?**)), in which we produce high quality results from the start without requiring a trained model. When a new input arrives, the system can choose to asynchronously query the crowd on *parts* of the input it is uncertain about (e.g. query about the label of a single token in a sentence). After collecting enough evidence the system makes a prediction. The goal is to maintain high accuracy by initially using the crowd as a crutch, but gradually becoming more self-sufficient as the model improves. Online learning (Cesa-Bianchi and Lugosi, 2006) and online active learning (Helmbold and Panizza, 1997; Sculley, 2007; Chu et al., 2011) are different in that they do not actively seek new information *prior* to making a prediction, and cannot maintain high accuracy independent of the number of data instances seen so far. Active classification (Gao and Koller, 2011), like us, strategically seeks information (by querying a subset of labels) prior to prediction, but it is based on a static policy, whereas we improve the model during test time based on observed data.

To determine which queries to make, we model on-the-job learning as a stochastic game based on a CRF prediction model. We use Bayesian decision theory to tradeoff latency, cost, and accuracy in a principled manner. Our framework naturally gives rise to intuitive strategies: To achieve high accuracy, we should ask for redundant labels to offset the noisy responses. To achieve low latency, we should issue queries in parallel, whereas if latency is unimportant, we should issue queries sequentially in order to be more adaptive. Computing the optimal policy is intractable, so we develop an approximation based on Monte Carlo tree search (Kocsis and Szepesvári, 2006) and progressive widening to reason

Figure 5.1: Named entity recognition on tweets in on-the-job learning.

about continuous time (Coulom, 2007).

We implemented and evaluated our system on three different tasks: named-entity recognition, sentiment classification, and image classification. On the NER task we obtained more than an order of magnitude reduction in cost compared to full human annotation, while boosting performance relative to the expert provided labels. We also achieve a 8% F1 improvement over having a single human label the whole set, and a 28% F1 improvement over online learning. An open-source implementation of our system, dubbed LENSE for "Learning from Expensive Noisy Slow Experts" is available at http://www.github.com/keenon/lense.

## 5.2 Problem formulation

Consider a structured prediction problem from input $\mathbf{x} = (x_1, \ldots, x_n)$ to output $\mathbf{y} = (y_1, \ldots, y_n)$. For example, for named-entity recognition (NER) on tweets, $\mathbf{x}$ is a sequence of words in the tweet (e.g., "*on George str.*") and $\mathbf{y}$ is the corresponding sequence of labels (e.g., NONE LOCATION LOCATION). The full set of labels of PERSON, LOCATION, RESOURCE, and NONE.

In the *on-the-job learning* setting, inputs arrive in a stream. On each input $\mathbf{x}$, we make zero or more queries $q_1, q_2, \ldots$ on the crowd to obtain labels (potentially more than once) for any positions in $\mathbf{x}$. The responses $r_1, r_2, \ldots$ come back asynchronously, which are

incorporated into our current prediction model $p_\theta$. Figure 5.2 (left) shows one possible outcome: We query positions $q_1 = 2$ ("*George*") and $q_2 = 3$ ("*str.*"). The first query returns $r_1 = \text{LOCATION}$, upon which we make another query on the the same position $q_3 = 3$ ("*George*"), and so on. When we have sufficient confidence about the entire output, we return the most likely prediction $\hat{y}$ under the model. Each query $q_i$ is issued at time $s_i$ and the response comes back at time $t_i$. Assume that each query costs $m$ cents. Our goal is to choose queries to maximize accuracy, minimize latency and cost.

We make several remarks about this setting: First, we must make a prediction $\hat{y}$ on each input $x$ in the stream, unlike in active learning, where we are only interested in the pool or stream of examples for the purposes of building a good model. Second, the responses are used to update the prediction model, like in online learning. This allows the number of queries needed (and thus cost and latency) to decrease over time without compromising accuracy.

## 5.3  Model

We model on-the-job learning as a stochastic game with two players: the system and the crowd. The game starts with the system receiving input $x$ and ends when the system turns in a set of labels $y = (y_1, \ldots, y_n)$. During the system's turn, the system may choose a query action $q \in \{1, \ldots, n\}$ to ask the crowd to label $y_q$. The system may also choose the wait action ($q = \varnothing_W$) to wait for the crowd to respond to a pending query or the return action ($q = \varnothing_R$) to terminate the game and return its prediction given responses received thus far. The system can make as many queries in a row (i.e. simultaneously) as it wants, before deciding to wait or turn in.[1] When the wait action is chosen, the turn switches to the crowd, which provides a response $r$ to one pending query, and advances the game clock by the time taken for the crowd to respond. The turn then immediately reverts back to the system. When the game ends (the system chooses the return action), the system evaluates a utility that depends on the accuracy of its prediction, the number of queries issued and the total time taken. The system should choose query and wait actions to maximize the utility of the prediction eventually returned.

---

[1] This rules out the possibility of launching a query midway through waiting for the next response. However, we feel like this is a reasonable limitation that significantly simplifies the search space.

(a) **Incorporating information from responses.** The bar graphs represent the marginals over the labels for each token (indicated by the first character) at different points in time. The two timelines show how the system updates its confidence over labels based on the crowd's responses. The system continues to issue queries until it has sufficient confidence on its labels. See the paragraph on behavior in Section 5.3 for more information.

(b) **Game tree.** An example of a partial game tree constructed by the system when deciding which action to take in the state $\sigma = (1, (3), (0), (\emptyset), (\emptyset))$, i.e. the query $q_1 = 3$ has already been issued and the system must decide whether to issue another query or wait for a response to $q_1$.

Figure 5.2: Example behavior while running structure prediction on the tweet "Soup on George str." We omit the RESOURCE from the game tree for visual clarity.

In the rest of this section, we describe the details of the game tree, our choice of utility and specify models for crowd responses, followed by a brief exploration of behavior admitted by our model.

**Game tree.** Let us now formalize the game tree in terms of its states, actions, transitions and rewards; see Figure 5.2b for an example. The *game state* $\sigma = (t_{\text{now}}, \mathbf{q}, \mathbf{s}, \mathbf{r}, \mathbf{t})$ consists of the current time $t_{\text{now}}$, the actions $\mathbf{q} = (q_1, \ldots, q_{k-1})$ that have been issued at times $\mathbf{s} = (s_1, \ldots, s_{k-1})$ and the responses $\mathbf{r} = (r_1, \ldots, r_{k-1})$ that have been received at times $\mathbf{t} = (t_1, \ldots, t_{k-1})$. Let $r_j = \emptyset$ and $t_j = \emptyset$ iff $q_j$ is not a query action or its responses have not been received by time $t_{\text{now}}$.

During the system's turn, when the system chooses an action $q_k$, the state is updated to $\sigma' = (t_{\text{now}}, \mathbf{q}', \mathbf{s}', \mathbf{r}', \mathbf{t}')$, where $\mathbf{q}' = (q_1, \ldots, q_k)$, $\mathbf{s}' = (s_1, \ldots, s_{k-1}, t_{\text{now}})$, $\mathbf{r}' = (r_1, \ldots, r_{k-1}, \emptyset)$ and $\mathbf{t}' = (t_1, \ldots, t_{k-1}, \emptyset)$. If $q_k \in \{1, \ldots n\}$, then the system chooses another action from the new state $\sigma'$. If $q_k = \emptyset_W$, the crowd makes a stochastic move from $\sigma'$. Finally, if $q_k = \emptyset_R$, the game ends, and the system returns its best estimate of the labels using the responses it has received and obtains a utility $U(\sigma)$ (defined later).

Let $F = \{1 \le j \le k - 1 \mid q_j \ne \emptyset_W \wedge r_j = \emptyset\}$ be the set of *in-flight* requests. During the crowd's turn (i.e. after the system chooses $\emptyset_W$), the next response from the crowd, $j^* \in F$, is chosen: $j^* = \arg\min_{j \in F} t'_j$ where $t'_j$ is sampled from the *response-time model*, $t'_j \sim p_{\text{T}}(t'_j \mid s_j, t'_j > t_{\text{now}})$, for each $j \in F$. Finally, a response is sampled using a response model, $r'_{j^*} \sim p(r'_{j^*} \mid \mathbf{x}, \mathbf{r})$, and the state is updated to $\sigma' = (t_{j^*}, \mathbf{q}, \mathbf{s}, \mathbf{r}', \mathbf{t}')$, where $\mathbf{r}' = (r_1, \ldots, r'_{j^*}, \ldots, r_k)$ and $\mathbf{t}' = (t_1, \ldots, t'_{j^*}, \ldots, t_k)$.

**Utility.** Under Bayesian decision theory, the *optimal choice* for an action in state $\sigma = (t_{\text{now}}, \mathbf{q}, \mathbf{r}, \mathbf{s}, \mathbf{t})$ is the one that attains the maximum expected utility (i.e. value) for the game starting at $\sigma$. Recall that the system can return at any time, at which point it receives a utility that trades off two things: The first is the accuracy of the MAP estimate according to the model's best guess of $\mathbf{y}$ incorporating all responses received by time $\tau$. The second is the cost of making queries: a (monetary) cost $w_{\text{M}}$ per query made and penalty of $w_{\text{T}}$ per

unit of time taken. Formally, we define the utility to be:

$$U(\sigma) \stackrel{\text{def}}{=} \text{ExpAcc}(p(\mathbf{y} \mid \mathbf{x}, \mathbf{q}, \mathbf{s}, \mathbf{r}, \mathbf{t})) - (n_{\text{Q}} w_{\text{M}} + t_{\text{now}} w_{\text{T}}), \tag{5.1}$$

$$\text{ExpAcc}(p) = \mathbb{E}_{p(\mathbf{y})}[\text{Accuracy}(\arg \max_{\mathbf{y}'} p(\mathbf{y}'))], \tag{5.2}$$

where $n_{\text{Q}} = |\{j \mid q_j \in \{1, \ldots, n\}|$ is the number of queries made, $p(\mathbf{y} \mid \mathbf{x}, \mathbf{q}, \mathbf{s}, \mathbf{r}, \mathbf{t})$ is a prediction model that incorporates the crowd's responses.

The utility of wait and return actions is computed by taking expectations over subsequent trajectories in the game tree. This is intractable to compute exactly, so we propose an approximate algorithm in Section 5.4.

**Environment model.** The final component is a model of the environment (crowd). Given input $\mathbf{x}$ and queries $\mathbf{q} = (q_1, \ldots, q_k)$ issued at times $\mathbf{s} = (s_1, \ldots, s_k)$, we define a distribution over the output $\mathbf{y}$, responses $\mathbf{r} = (r_1, \ldots, r_k)$ and response times $\mathbf{t} = (t_1, \ldots, t_k)$ as follows:

$$p(\mathbf{y}, \mathbf{r}, \mathbf{t} \mid \mathbf{x}, \mathbf{q}, \mathbf{s}) \stackrel{\text{def}}{=} (() \mathbf{y} \mid \mathbf{x}) \prod_{i=1}^{k} p_{\text{R}}(r_i \mid y_{q_i}) p_{\text{T}}(t_i \mid s_i). \tag{5.3}$$

The three components are as follows: $(() \mathbf{y} \mid \mathbf{x})$ is the *prediction model* (e.g. a standard linear-chain CRF); $p_{\text{R}}(r \mid y_q)$ is the *response model* which describes the distribution of the crowd's response $r$ for a given a query $q$ when the true answer is $y_q$; and $p_{\text{T}}(t_i \mid s_i)$ specifies the latency of query $q_i$. The CRF model $(() \mathbf{y} \mid \mathbf{x})$ is learned based on all actual responses (not simulated ones) using AdaGrad. To model annotation errors, we set $p_{\text{R}}(r \mid y_q) = 0.7$ iff $r = y_q$,[2] and distribute the remaining probability for $r$ uniformly. Given this full model, we can compute $p(r' \mid \mathbf{x}, \mathbf{r}, q)$ simply by marginalizing out $\mathbf{y}$ and $\mathbf{t}$ from (5.3). When conditioning on $\mathbf{r}$, we ignore responses that have not yet been received (i.e. when $r_j = \emptyset$ for some $j$).

---

[2] We found the humans we hired were roughly 70% accurate in our experiments

**Behavior.** Let's look at typical behavior that we expect the model and utility to capture. Figure 5.2a shows how the marginals over the labels change as the crowd provides responses for our running example, i.e. named entity recognition for the sentence "Soup on George str.". In the both timelines, the system issues queries on "Soup" and "George" because it is not confident about its predictions for these tokens. In the first timeline, the crowd correctly responds that "Soup" is a resource and that "George" is a location. Integrating these responses, the system is also more confident about its prediction on "str.", and turns in the correct sequence of labels. In the second timeline, a crowd worker makes an error and labels "George" to be a person. The system still has uncertainty on "George" and issues an additional query which receives a correct response, following which the system turns in the correct sequence of labels. While the answer is still correct, the system could have taken less time to respond by making an additional query on "George" at the very beginning.

## 5.4 Game playing

In Section 5.3 we modeled on-the-job learning as a stochastic game played between the system and the crowd. We now turn to the problem of actually finding a policy that maximizes the expected utility, which is, of course, intractable because of the large state space.

Our algorithm (Algorithm **??**) combines ideas from Monte Carlo tree search (Kocsis and Szepesvári, 2006) to systematically explore the state space and progressive widening (Coulom, 2007) to deal with the challenge of continuous variables (time). Some intuition about the algorithm is provided below. When simulating the system's turn, the next state (and hence action) is chosen using the upper confidence tree (UCT) decision rule that trades off maximizing the value of the next state (exploitation) with the number of visits (exploration). The crowd's turn is simulated based on transitions defined in Section 5.3. To handle the unbounded fanout during the crowd's turn, we use progressive widening that maintains a current set of "active" or "explored" states, which is gradually grown with time. Let $N(\sigma)$ be the number of times a state has been visited, and $C(\sigma)$ be all successor states that the algorithm has sampled.

---

**Algorithm 2** Approximating expected utility with MCTS and progressive widening

---

1: For all $\sigma$, $N(\sigma) \leftarrow 0$, $V(\sigma) \leftarrow 0$, $C(\sigma) \leftarrow [\,]$    ▷ Initialize visits, utility sum, and children
2: **function** MONTECARLOVALUE(state $\sigma$)
3:   increment $N(\sigma)$
4:   **if** system's turn **then**
5:    $\sigma' \leftarrow \arg\max_{\sigma'} \left\{ \frac{V(\sigma')}{N(\sigma')} + c\sqrt{\frac{\log N(\sigma)}{N(\sigma')}} \right\}$   ▷ Choose next state $\sigma'$ using UCT
6:    $v \leftarrow$ MONTECARLOVALUE($\sigma'$)
7:    $V(\sigma) \leftarrow V(\sigma) + v$        ▷ Record observed utility
8:    **return** $v$
9:   **else if** crowd's turn **then**
10:    **if** $\max(1, \sqrt{N(\sigma)}) \leq |C(\sigma)|$ **then**   ▷ Restrict continuous samples using PW
11:     $\sigma'$ is sampled from set of already visited $C(\sigma)$ based on (5.3)
12:    **else**
13:     $\sigma'$ is drawn based on (5.3)
14:     $C(\sigma) \leftarrow C(\sigma) \cup \{[\sigma']\}$
15:    **end if**
16:    **return** MONTECARLOVALUE($\sigma'$)
17:   **else if** game terminated **then**
18:    **return** utility $U$ of $\sigma$ according to (5.1)
19:   **end if**
20: **end function**

---

## 5.5 Experiments

In this section, we empirically evaluate our approach on three tasks. While the on-the-job setting we propose is targeted at scenarios where there is no data to begin with, we use existing labeled datasets (Table 4.2) to have a gold standard.

**Baselines.**    We evaluated the following four methods on each dataset:

1. **Human n-query**: The majority vote of $n$ human crowd workers was used as a prediction.

2. **Online learning**: Uses a classifier that trains on the gold output for all examples seen so far and then returns the MLE as a prediction. This is the best possible offline system: it sees perfect information about all the data seen so far, but can not query the crowd while making a prediction.

3. **Threshold baseline**: Uses the following heuristic: For each label, $y_i$, we ask for $m$ queries such that $(1 - (() y_i \mid \mathbf{x})) \times 0.3^m \geq 0.98$. Instead of computing the expected marginals over the responses to queries in flight, we simply count the in-flight requests for a given variable, and reduces the uncertainty on that variable by a factor of $0.3$. The system continues launching requests until the threshold (adjusted by number of queries in flight) is crossed. Predictions are made using MLE on the model given responses. The baseline does not reason about time and makes all its queries at the very beginning.

4. **LENSE:** Our full system as described in Section 5.3.

**Implementation and crowdsourcing setup.**    We implemented the retainer model of Bernstein et al. (2011) on Amazon Mechanical Turk to create a "pool" of crowd workers that could respond to queries in real-time. The workers were given a short tutorial on each task before joining the pool to minimize systematic errors caused by misunderstanding the task.

---

[3]http://www.cnts.ua.ac.be/conll2003/ner/

[4] The original also includes a fifth tag for miscellaneous, however the definition for miscellaneos is complex, making it very difficult for non-expert crowd workers to provide accurate labels.

| Dataset (Examples) | Task and notes | Features |
|---|---|---|
| **NER (657)** | We evaluate on the CoNLL-2003 NER task[3], a sequence labeling problem over English sentences. We only consider the four tags corresponding to persons, locations, organizations or none[4]. | We used standard features Finkel et al. (2005): the current word, current lemma, previous and next lemmas, lemmas in a window of size three to the left and right, word shape and word prefix and suffixes, as well as word embeddings. |
| **Sentiment (1800)** | We evaluate on a subset of the IMDB sentiment dataset **?** that consists of 2000 polar movie reviews; the goal is binary classification of documents into classes POS and NEG. | We used two feature sets, the first (UNIGRAMS) containing only word unigrams, and the second (RNN) that also contains sentence vector embeddings from Socher et al. (2013). |
| **Face (1784)** | We evaluate on a celebrity face classification task **?**. Each image must be labeled as one of the following four choices: Andersen Cooper, Daniel Craig, Scarlet Johansson or Miley Cyrus. | We used the last layer of a 11-layer AlexNet Krizhevsky et al. (2012) trained on ImageNet as input feature embeddings, though we leave back-propagating into the net to future work. |

Table 5.1: Datasets used in this paper and number of examples we evaluate on.

We paid workers $1.00 to join the retainer pool and an additional $0.01 per query (for NER, since response times were much faster, we paid $0.005 per query). Worker response times were generally in the range of 0.5–2 seconds for NER, 10–15 seconds for Sentiment, and 1–4 seconds for Faces.

When running experiments, we found that the results varied based on the current worker quality. To control for variance in worker quality across our evaluations of the different methods, we collected 5 worker responses and their delays on each label ahead of time[5]. During simulation we sample the worker responses and delays without replacement from this frozen pool of worker responses.

---

[5]These datasets are available in the code repository for this paper

| | Named Entity Recognition | | | | | | Face Identification | | |
|---|---|---|---|---|---|---|---|---|---|
| **System** | **Delay/tok** | **Qs/tok** | **PER $F_1$** | **LOC $F_1$** | **ORG $F_1$** | **$F_1$** | **Latency** | **Qs/ex** | **Acc.** |
| 1-vote | 467 ms | 1.0 | 90.2 | 78.8 | 71.5 | 80.2 | 1216 ms | 1.0 | 93.6 |
| 3-vote | 750 ms | 3.0 | 93.6 | 85.1 | 74.5 | 85.4 | 1782 ms | 3.0 | 99.1 |
| 5-vote | 1350 ms | 5.0 | **95.5** | 87.7 | 78.7 | 87.3 | 2103 ms | 5.0 | 99.8 |
| Online | n/a | n/a | 56.9 | 74.6 | 51.4 | 60.9 | n/a | n/a | 79.9 |
| Threshold | 414 ms | 0.61 | 95.2 | **89.8** | 79.8 | 88.3 | 1680 ms | 2.66 | 93.5 |
| **LENSE** | **267 ms** | **0.45** | 95.2 | 89.7 | **81.7** | **88.8** | 1590 ms | 2.37 | 99.2 |

Table 5.2: Results on NER and Face tasks comparing latencies, queries per token (Qs/tok) and performance metrics ($F_1$ for NER and accuracy for Face).



Figure 5.4: Comparing $F_1$ and queries per token on the NER task over time. The left graph compares LENSE to online learning (which cannot query humans at test time). This highlights that LENSE maintains high $F_1$ scores even with very small training set sizes, by falling back the crowd when it is unsure. The right graph compares query rate over time to 1-vote. This clearly shows that as the model learns, it needs to query the crowd less.

**Summary of results.** Table **??** and Table **??** summarize the performance of the methods on the three tasks. On all three datasets, we found that on-the-job learning outperforms machine and human-only comparisons on both quality and cost. On NER, we achieve an $F_1$ of $88.4\%$ at more than an order of magnitude reduction on the cost of achieving comporable quality result using the 5-vote approach. On Sentiment and Faces, we reduce costs for a comparable accuracy by a factor of around 2. For the latter two tasks, both on-the-job learning methods perform less well than in NER. We suspect this is due to the presence of a dominant class ("none") in NER that the model can very quickly learn to expend almost no effort on. LENSE outperforms the threshold baseline, supporting the importance of Bayesian decision theory.

Figure 5.4 tracks the performance and cost of LENSE over time on the NER task.

Figure 5.3: Queries per example for LENSE on Sentiment. With simple UNIGRAM features, the model quickly learns it does not have the capacity to answer confidently and must query the crowd. With more complex RNN features, the model learns to be more confident and queries the crowd less over time.

| System | Latency | Qs/ex | Acc. |
|---|---|---|---|
| 1-vote | 6.6 s | 1.00 | 89.2 |
| 3-vote | 10.9 s | 3.00 | 95.8 |
| 5-vote | 13.5 s | 5.00 | 98.7 |
| UNIGRAMS | | | |
| Online | n/a | n/a | 78.1 |
| Threshold | 10.9 s | 2.99 | 95.9 |
| **LENSE** | 11.7 s | 3.48 | 98.6 |
| RNN | | | |
| Online | n/a | n/a | 85.0 |
| Threshold | 11.0 s | 2.85 | 96.0 |
| **LENSE** | 11.0 s | 3.19 | 98.6 |

Table 5.3: Results on the Sentiment task comparing latency, queries per example and accuracy.

LENSE is not only able to consistently outperform other baselines, but the cost of the system steadily reduces over time. On the NER task, we find that LENSE is able to trade off time to produce more accurate results than the 1-vote baseline with fewer queries by waiting for responses before making another query.

While on-the-job learning allows us to deploy quickly and ensure good results, we would like to eventually operate without crowd supervision. Figure 5.3, we show the number of queries per example on Sentiment with two different features sets, UNIGRAMS and RNN (as described in Table 4.2). With simpler features (UNIGRAMS), the model saturates early and we will continue to need to query to the crowd to achieve our accuracy target (as specified by the loss function). On the other hand, using richer features (RNN) the model is able to learn from the crowd and the amount of supervision needed reduces over time. Note that even when the model capacity is limited, LENSE is able to guarantee a consistent, high level of performance.

**Reproducibility.** All code, data, and experiments for this paper are available on CodaLab at https://www.codalab.org/worksheets/0x2ae89944846444539c2d08a0b7ff3f6f/.

## 5.6 Related Work

On-the-job learning draws ideas from many areas: online learning, active learning, active classification, crowdsourcing, and structured prediction.

**Online learning.** The fundamental premise of online learning is that algorithms should improve with time, and there is a rich body of work in this area (Cesa-Bianchi and Lugosi, 2006). In our setting, algorithms not only improve over time, but maintain high accuracy from the beginning, whereas regret bounds only achieve this asymptotically.

**Active learning.** Active learning (see Settles (2010) for a survey) algorithms strategically select most informative examples to build a classifier. Online active learning (Helmbold and Panizza, 1997; Sculley, 2007; Chu et al., 2011) performs active learning in the online setting. Several authors have also considered using crowd workers as a noisy oracle e.g. (Donmez and Carbonell, 2008; Golovin et al., 2010). It differs from our setup in that it assumes that labels can only be observed *after* classification, which makes it nearly impossible to maintain high accuracy in the beginning.

**Active classification.** Active classification Greiner et al. (2002); Chai et al. (2004); Esmeir and Markovitch (2007) asks what are the *most informative features* to measure at test time. Existing active classification algorithms rely on having a fully labeled dataset which is used to learn a static policy for when certain features should be queried, which does not change at test time. On-the-job learning differs from active classification in two respects: true labels are *never* observed, and our system improves itself at test time by learning a stronger model. A notable exception is Legion:AR **?**, which like us operates in on-the-job learning setting to for real-time activity classification. However, they do not explore the machine learning foundations associated with operating in this setting, which is the aim of this paper.

**Crowdsourcing.** A burgenoning subset of the crowdsourcing community overlaps with machine learning. One example is *Flock* Cheng and Bernstein (2015), which first crowdsources the identification of features for an image classification task, and then asks the

crowd to annotate these features so it can learn a decision tree. In another line of work, *TurKontrol* Dai et al. (2010) models individual crowd worker reliability to optimize the number of human votes needed to achieve confident consensus using a POMDP.

**Structured prediction.** An important aspect our prediction tasks is that the output is structured, which leads to a much richer setting for one-the-job learning. Since tags are correlated, the importance of a coherent framework for optimizing querying resources is increased. Making active partial observations on structures and has been explored in the measurements framework of Liang et al. (2009) and in the distant supervision setting Angeli et al. (2014).

## 5.7   Conclusion

We have introduced a new framework that learns from (noisy) crowds *on-the-job* to maintain high accuracy, and reducing cost significantly over time. The technical core of our approach is modeling the on-the-job setting as a stochastic game and using ideas from game playing to approximate the optimal policy. We have built a system, LENSE, which obtains significant cost reductions over a pure crowd approach and significant accuracy improvements over a pure ML approach.

# Chapter 6

# Discussion: where do we go from here?

placeholder.

Linguistic notions of structure. Thematic progression. Bad at cohesion.

# Bibliography

H. Adel, B. Roth, and H. Schütze. 2016. Comparing convolutional neural networks to traditional models for slot filling. In *Human Language Technology and North American Association for Computational Linguistics (HLT/NAACL)*.

G. Angeli, J. Tibshirani, J. Y. Wu, and C. D. Manning. 2014. Combining distant and partial supervision for relation extraction. In *Empirical Methods in Natural Language Processing (EMNLP)*.

J. A. Aslam, V. Pavlu, and E. Yilmaz. 2006. A statistical method for system evaluation using incomplete judgments. In *ACM Special Interest Group on Information Retreival (SIGIR)*, pages 541–548.

J. Berant, A. Chou, R. Frostig, and P. Liang. 2013. Semantic parsing on Freebase from question-answer pairs. In *Empirical Methods in Natural Language Processing (EMNLP)*.

M. S. Bernstein, J. Brandt, R. C. Miller, and D. R. Karger. 2011. Crowds in two seconds: Enabling realtime crowd-powered interfaces. In *User Interface Software and Technology*, pages 33–42.

M. S. Bernstein, G. Little, R. C. Miller, B. Hartmann, M. S. Ackerman, D. R. Karger, D. Crowell, and K. Panovich. 2010. Soylent: a word processor with a crowd inside. In *Symposium on User Interface Software and Technology*, pages 313–322.

R. Brandow, K. Mitze, and L. F. Rau. 1995. Automatic condensation of electronic publications by sentence selection. *Information Processing and Management*, 31:675–685.

C. Buckley, D. Dimmick, I. Soboroff, and E. Voorhees. 2007. Bias and the limits of pooling for large collections. In *ACM Special Interest Group on Information Retreival (SIGIR)*.

C. Buckley and E. M. Voorhees. 2004. Retrieval evaluation with incomplete information. In *ACM Special Interest Group on Information Retreival (SIGIR)*, pages 25–32.

R. L. Burden and J. D. Faires. 1985. *Numerical Analysis (3rd ed.)*. PWS Publishers.

N. Cesa-Bianchi and G. Lugosi. 2006. *Prediction, learning, and games*. Cambridge University Press.

A. Chaganty, A. Paranjape, P. Liang, and C. Manning. 2017. Importance sampling for unbiased on-demand evaluation of knowledge base population. In *Empirical Methods in Natural Language Processing (EMNLP)*.

X. Chai, L. Deng, Q. Yang, and C. X. Ling. 2004. Test-cost sensitive naive Bayes classification. In *International Conference on Data Mining*, pages 51–58.

C. Chang, R. Yang, L. Chen, X. Zhou, and K. Yu. 2017. Affordable on-line dialogue policy learning. In *Empirical Methods in Natural Language Processing (EMNLP)*, pages 223–231.

J. Cheng and M. S. Bernstein. 2015. Flock: Hybrid Crowd-Machine learning classifiers. In *Proceedings of the 18th ACM Conference on Computer Supported Cooperative Work & Social Computing*, pages 600–611.

W. Chu, M. Zinkevich, L. Li, A. Thomas, and B. Tseng. 2011. Unbiased online active learning in data streams. In *International Conference on Knowledge Discovery and Data Mining (KDD)*, pages 195–203.

J. M. Conroy and H. T. Dang. 2008. Mind the gap : Dangers of divorcing evaluations of summary content from linguistic quality. In *International Conference on Computational Linguistics (COLING)*, pages 145–152.

G. V. Cormack, C. R. Palmer, and C. L. A. Clarke. 1998. Efficient construction of large test collections. In *ACM Special Interest Group on Information Retreival (SIGIR)*.

R. Coulom. 2007. Computing elo ratings of move patterns in the game of go. *Computer Games Workshop*.

P. Dai, Mausam, and D. S. Weld. 2010. Decision-theoretic control of crowd-sourced workflows. In *Association for the Advancement of Artificial Intelligence (AAAI)*.

H. T. Dang. 2006. Overview of DUC 2006. In *Document Understanding Conference*.

H. T. Dang. 2016. Cold start knowledge base population at TAC KBP 2016. *Text Analytics Conference*.

J. Deng, W. Dong, R. Socher, L. Li, K. Li, and L. Fei-Fei. 2009. ImageNet: A large-scale hierarchical image database. In *Computer Vision and Pattern Recognition (CVPR)*, pages 248–255.

M. Denkowski and A. Lavie. 2014. Meteor universal: Language specific translation evaluation for any target language. In *Workshop on Statistical Machine Translation*.

P. Donmez and J. G. Carbonell. 2008. Proactive learning: cost-sensitive active learning with multiple imperfect oracles. In *Conference on Information and Knowledge Management (CIKM)*, pages 619–628.

O. Dusek, J. Novikova, and V. Rieser. 2017. Referenceless quality estimation for natural language generation. *arXiv*.

J. Ellis, X. Li, K. Griffitt, and S. M. Strassel. 2012. Linguistic resources for 2012 knowledge base population evaluations. *Text Analytics Conference*.

S. Esmeir and S. Markovitch. 2007. Anytime induction of cost-sensitive trees. In *Advances in Neural Information Processing Systems (NIPS)*, pages 425–432.

A. Fader, L. Zettlemoyer, and O. Etzioni. 2014. Open question answering over curated and extracted knowledge bases. In *International Conference on Knowledge Discovery and Data Mining (KDD)*, pages 1156–1165.

J. R. Finkel, T. Grenager, and C. Manning. 2005. Incorporating non-local information into information extraction systems by Gibbs sampling. In *Association for Computational Linguistics (ACL)*, pages 363–370.

T. Gao and D. Koller. 2011. Active classification based on value of classifier. In *Advances in Neural Information Processing Systems (NIPS)*, pages 1062–1070.

D. Golovin, A. Krause, and D. Ray. 2010. Near-optimal Bayesian active learning with noisy observations. In *Advances in Neural Information Processing Systems (NIPS)*, pages 766–774.

E. Greensmith, P. L. Bartlett, and J. Baxter. 2004. Variance reduction techniques for gradient estimates in reinforcement learning. *Journal of Machine Learning Research (JMLR)*, 5:1471–1530.

R. Greiner, A. J. Grove, and D. Roth. 2002. Learning cost-sensitive active classifiers. *Artificial Intelligence*, 139(2):137–174.

S. Han, J. Bang, S. Ryu, and G. G. Lee. 2015. Exploiting knowledge base to generate responses for natural language dialog listening agents. *16th Annual Meeting of the Special Interest Group on Discourse and Dialogue*, pages 129–133.

D. K. Harman. 1993. The first text retrieval conference (trec-1) rockville, md, u.s.a., 4-6 november, 1992. *Information Processing and Management*, 29:411–414.

L. He, M. Lewis, and L. Zettlemoyer. 2015. Question-answer driven semantic role labeling: Using natural language to annotate natural language. In *Empirical Methods in Natural Language Processing (EMNLP)*.

D. Helmbold and S. Panizza. 1997. Some label efficient learning results. In *Conference on Learning Theory (COLT)*, pages 218–230.

K. M. Hermann, T. Kočiský, E. Grefenstette, L. Espeholt, W. Kay, M. Suleyman, and P. Blunsom. 2015. Teaching machines to read and comprehend. In *Advances in Neural Information Processing Systems (NIPS)*.

H. Ji, R. Grishman, and H. Trang Dang. 2011. Overview of the TAC 2011 knowledge base population track. In *Text Analytics Conference*.

K. S. Jones and C. V. Rijsbergen. 1975. Report on the need for and provision of an "ideal test collection. *Information Retrieval Test Collection*.

A. Kalyanpur, B. K. Boguraev, S. Patwardhan, J. W. Murdock, A. Lally, C. A. Welty, J. M. Prager, B. Coppola, A. Fokoue-Nkoutche, L. Zhang, Y. Pan, and Z. M. Qui. 2012. Structured data and inference in deepqa. *IBM Journal of Research and Development*, 56:351–364.

T. Kočisky, J. Schwarz, P. Blunsom, C. Dyer, K. M. Hermann, G. Melis, and E. Grefenstette. 2017. The NarrativeQA reading comprehension challenge. *arXiv preprint arXiv:1712.07040*.

L. Kocsis and C. Szepesvári. 2006. Bandit based Monte-Carlo planning. In *European Conference on Machine Learning (ECML)*, pages 282–293.

N. Kokkalis, T. Köhn, C. Pfeiffer, D. Chornyi, M. S. Bernstein, and S. R. Klemmer. 2013. Emailvalet: Managing email overload through private, accountable crowdsourcing. In *Conference on Computer Supported Cooperative Work*, pages 1291–1300.

A. Krizhevsky, I. Sutskever, and G. E. Hinton. 2012. Imagenet classification with deep convolutional neural networks. In *Advances in Neural Information Processing Systems (NIPS)*, pages 1097–1105.

A. Lavie and M. Denkowski. 2009. The meteor metric for automatic evaluation of machine translation. *Machine Translation*, 23.

C. Li, J. Weng, Q. He, Y. Yao, A. Datta, A. Sun, and B. Lee. 2012. Twiner: named entity recognition in targeted twitter stream. In *ACM Special Interest Group on Information Retreival (SIGIR)*, pages 721–730.

P. Liang, M. I. Jordan, and D. Klein. 2009. Learning from measurements in exponential families. In *International Conference on Machine Learning (ICML)*.

C. Lin and M. Rey. 2004. Looking for a few good metrics: ROUGE and its evaluation. In *NTCIR Workshop*.

T. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Doll'ar, and C. L. Zitnick. 2014. Microsoft COCO: Common objects in context. In *European Conference on Computer Vision (ECCV)*, pages 740–755.

A. Liu, S. Soderland, J. Bragg, C. H. Lin, X. Ling, and D. S. Weld. 2016a. Effective crowd annotation for relation extraction. In *North American Association for Computational Linguistics (NAACL)*, pages 897–906.

C. Liu, R. Lowe, I. V. Serban, M. Noseworthy, L. Charlin, and J. Pineau. 2016b. How NOT to evaluate your dialogue system: An empirical study of unsupervised evaluation metrics for dialogue response generation. In *Empirical Methods in Natural Language Processing (EMNLP)*.

R. Lowe, M. Noseworthy, I. V. Serban, N. Angelard-Gontier, Y. Bengio, and J. Pineau. 2017a. Towards an automatic turing test: Learning to evaluate dialogue responses. In *Association for Computational Linguistics (ACL)*.

R. T. Lowe, N. Pow, I. Serban, L. Charlin, C. Liu, and J. Pineau. 2017b. Training end-to-end dialogue systems with the ubuntu dialogue corpus. *Dialogue and Discourse*, 8.

H. P. Luhn. 1958. The automatic creation of literature abstracts. *IBM Journal of Research and Development*, 2:159–165.

I. Mani, G. Klein, L. Hirschman, T. Firmin, D. House, and B. Sundheim. 1999. The TIP-STER SUMMAC text summarization evaluation. In *European Association for Computational Linguistics (EACL)*.

C. D. Manning, M. Surdeanu, J. Bauer, J. Finkel, S. J. Bethard, and D. McClosky. 2014. The stanford coreNLP natural language processing toolkit. In *ACL system demonstrations*.

K. Mckeown, R. J. Passonneau, D. K. Elson, and J. Hirschberg. 2005. Do summaries help? a task-based evaluation of multi-document summarization. In *ACM Special Interest Group on Information Retreival (SIGIR)*.

V. Mnih, C. Szepesv'ari, and J. Audibert. 2008. Empirical berstein stopping. In *International Conference on Machine Learning (ICML)*.

R. Nallapati, B. Zhou, C. Gulcehre, B. Xiang, et al. 2016. Abstractive text summarization using sequence-to-sequence rnns and beyond. *arXiv preprint arXiv:1602.06023*.

T. Nguyen, M. Rosenberg, X. Song, J. Gao, S. Tiwary, R. Majumder, and L. Deng. 2016. MS MARCO: A human generated machine reading comprehension dataset. In *Workshop on Cognitive Computing at NIPS*.

J. Novikova, O. Dušek, A. C. Curry, and V. Rieser. 2017. Why we need new evaluation metrics for NLG. In *Empirical Methods in Natural Language Processing (EMNLP)*.

A. B. Owen. 2013. *Monte Carlo theory, methods and examples*.

M. Pagliardini, P. Gupta, and M. Jaggi. 2017. Unsupervised learning of sentence embeddings using compositional n-gram features. *arXiv*.

J. Paisley, D. M. Blei, and M. I. Jordan. 2012. Variational Bayesian inference with stochastic search. In *International Conference on Machine Learning (ICML)*, pages 1363–1370.

K. Papineni, S. Roukos, T. Ward, and W. Zhu. 2002. BLEU: A method for automatic evaluation of machine translation. In *Association for Computational Linguistics (ACL)*.

R. J. Passonneau and B. Carpenter. 2014. The benefits of a model of annotation. In *Association for Computational Linguistics (ACL)*.

R. J. Passonneau, A. Nenkova, K. McKeown, and S. Sigelman. 2005. Applying the pyramid method in DUC 2005. In *Document Understanding Conference*.

R. Paulus, C. Xiong, and R. Socher. 2018. A deep reinforced model for abstractive summarization. In *International Conference on Learning Representations (ICLR)*.

E. Pavlick, H. Ji, X. Pan, and C. Callison-Burch. 2016. The gun violence database: A new task and data set for NLP. In *Empirical Methods in Natural Language Processing (EMNLP)*, pages 1018–1024.

R. Ranganath, S. Gerrish, and D. Blei. 2014. Black box variational inference. In *Artificial Intelligence and Statistics (AISTATS)*, pages 814–822.

L. Ratinov, D. Roth, D. Downey, and M. Anderson. 2011. Local and global algorithms for disambiguation to Wikipedia. In *Association for Computational Linguistics (ACL)*.

S. Reddy, M. Lapata, and M. Steedman. 2014. Large-scale semantic parsing without question-answer pairs. *Transactions of the Association for Computational Linguistics (TACL)*, 2(10):377–392.

B. D. Ripley. 2009. *Stochastic simulation*. John Wiley & Sons.

T. Sakai and N. Kando. 2008. On information retrieval metrics designed for evaluation with incomplete relevance assessments. In *ACM Special Interest Group on Information Retreival (SIGIR)*, pages 447–470.

D. Sculley. 2007. Online active learning methods for fast label-efficient spam filtering. In *Conference on Email and Anti-spam (CEAS)*.

A. See, P. J. Liu, and C. D. Manning. 2017. Get to the point: Summarization with pointer-generator networks. In *Association for Computational Linguistics (ACL)*.

B. Settles. 2010. Active learning literature survey. Technical report, University of Wisconsin, Madison.

M. Snover, B. Dorr, R. Schwartz, L. Micciulla, and J. Makhoul. 2006. A study of translation edit rate with targeted human annotation. In *Association for Machine Translation in the Americas*, pages 223–231.

R. Socher, A. Perelygin, J. Y. Wu, J. Chuang, C. D. Manning, A. Y. Ng, and C. Potts. 2013. Recursive deep models for semantic compositionality over a sentiment treebank. In *Empirical Methods in Natural Language Processing (EMNLP)*.

C. Tan, F. Wei, N. Yang, W. Lv, and M. Zhou. 2018. S-Net: From answer extraction to answer generation for machine reading comprehension. In *Association for the Advancement of Artificial Intelligence (AAAI)*.

D. Vannella, D. Jurgens, D. Scarfini, D. Toscani, and R. Navigli. 2014. Validating and extending semantic knowledge bases using video games with a purpose. In *Association for Computational Linguistics (ACL)*, pages 1294–1304.

R. Vedantam, C. L. Zitnick, and D. Parikh. 2015. CIDEr: Consensus-based image description evaluation. In *Computer Vision and Pattern Recognition (CVPR)*, pages 4566–4575.

W. E. Webber. 2010. *Measurement in Information Retrieval Evaluation*. Ph.D. thesis, University of Melbourne.

D. Weissenborn, G. Wiese, and L. Seiffe. 2017. Making neural QA as simple as possible but not simpler. In *Computational Natural Language Learning (CoNLL)*.

Y. Wu, M. Schuster, Z. Chen, Q. V. Le, M. Norouzi, W. Macherey, M. Krikun, Y. Cao, Q. Gao, K. Macherey, et al. 2016. Google's neural machine translation system: Bridging the gap between human and machine translation. *arXiv preprint arXiv:1609.08144*.

E. Yilmaz, E. Kanoulas, and J. A. Aslam. 2008. A simple and efficient sampling method for estimating AP and NDCG. In *ACM Special Interest Group on Information Retreival (SIGIR)*, pages 603–610.

J. Zobel. 1998. How reliable are the results of large-scale information retrieval experiments? In *ACM Special Interest Group on Information Retreival (SIGIR)*.